

# Model Monster's AI Glossary | Model Monster

---

**Reference copy.** Source: <https://modelmonster.ai/glossary/>

Retrieved 2026-06-05 · Process Mechanics PLLC / Van Lindberg · retained with permission for internal reference.

This glossary is a shared vocabulary for discussing AI technology, deployment, and related governance, procurement, and policy topics. Definitions are intended to be primarily descriptive, with some pointers as to why certain terms matter in the legal context. However, these definitions are limited and how a term applies in a given matter depends on the facts, jurisdiction, and the governing agreement.

## A

### A2A

See: [Agent-to-Agent Protocol](#)

### Abuse

C R

Misuse of an AI system or service (e.g., policy-violating content, fraud, phishing, scraping, automated harassment) or abuse of the underlying service (e.g., credential stuffing, model extraction attempts). “Abuse” is commonly used in safety policies, AUP enforcement, and monitoring programs.

See: [Acceptable Use Policy](#); [Content filtering](#); [Model extraction](#); [Rate limiting](#)

### Acceptable use

A general term for permitted use of a service. In many settings it refers to an Acceptable Use Policy (AUP) or similar contractual restrictions.

See: [Acceptable Use Policy](#); [Terms of Service](#)

### Acceptable Use Policy (AUP)

C

Rules (often incorporated into or referenced by Terms of Service) describing permitted and prohibited uses of an AI service and the provider’s enforcement options (e.g., warnings, suspension, termination). AUPs are commonly relevant to authorization questions, misuse response, and allocation of responsibility for prohibited uses. AUPs are frequently incorporated by reference into licenses.

See: [Content filtering](#); [Safety policy](#); [Terms of Service](#)

## Acceptance criteria

C

Defined, testable conditions for determining whether a deliverable or service meets requirements (e.g., performance thresholds, evaluation results, latency, groundedness, safety tests). In AI agreements, acceptance criteria often tie to a specific intended use and evaluation methodology.

See: [Benchmark](#); [Evaluation \(evals\)](#); [Intended use](#); [SLA/SLO](#); [Testing](#)

## Access control

C R

Technical and administrative mechanisms that limit who or what can access data, systems, and tools (e.g., authentication, authorization, role-based access control, least privilege). Access control is central to confidentiality, privacy, and security obligations.

See: [Authentication](#); [Authorization](#); [Connector](#); [Data security](#); [Least privilege](#)

## Accuracy

The proportion of predictions a model gets correct. Accuracy is context-dependent and often misleading without additional metrics: a model that predicts "no fraud" for every transaction achieves 99% accuracy if fraud occurs in only 1% of cases. Marketing claims citing accuracy should specify the dataset, task, and conditions; high accuracy on benchmarks may mask poor performance on important subgroups.

See: [Evaluation \(evals\)](#); [F1 score](#); [Precision](#); [Recall](#)

## Active learning

A machine learning approach where the model identifies which unlabeled examples would be most valuable to label, reducing annotation costs. Active learning strategies may involve human labelers viewing sensitive content, raising labor and content moderation considerations.

See: [Annotation](#); [Human evaluation](#); [Labeling](#)

## Adapter

A small set of parameters added to a frozen base model to customize behavior without full fine-tuning; includes techniques like LoRA. Adapters are often the deliverable in custom AI projects, so ownership, portability, and confidentiality terms are often stated explicitly. Frequently used as a way to provide organization-specific capabilities without creating full separate base models.

See: [Fine-tuning](#); [LoRA \(Low-Rank Adaptation\)](#); [Parameter-efficient fine-tuning](#)

## Adversarial attack

An attempt to cause a model to produce incorrect or harmful outputs through crafted inputs. Adversarial robustness is relevant to security representations, product liability, possibly-infringing intellectual property outputs, and contractual performance standards. This can be an attack against the AI system as a whole or even just a prompt designed to elicit an unwanted result.

See: [Evasion attack](#); [Jailbreak](#); [Prompt injection](#); [Red teaming](#)

## Adversarial example

A data instance purposefully perturbed to induce misclassification by a deployed model. Relevant to security testing, autonomous vehicle litigation, and content moderation failures.

See: [Adversarial attack](#); [Data poisoning](#); [Evaluation \(evals\)](#); [Robustness](#)

## Adversarial success

An AI system failure due to adversarial attack, where unwanted model output results in adverse effects like leakage of privileged data, violation of guardrails, expansion of privilege, or unwanted output. This defines the failure condition for many types of testing and is relevant to breach notification and incident response obligations.

See: [Data leakage](#); [Jailbreak](#); [Prompt injection](#)

## Adverse action

L R

A negative decision affecting an individual (commonly used in employment, credit, housing, insurance, and benefits contexts). AI-supported decisions that lead to adverse action may implicate notice, contestability, documentation, or anti-discrimination requirements depending on jurisdiction and sector.

See: [Algorithmic discrimination](#); [Automated decision-making](#); [Disparate impact](#)

## Agent / AI agent

C R

An embedded or dedicated software component that interacts with a person or system. An agent may or may not be "agentic" - i.e., it may or may not include autonomous or semi-autonomous decision-making or tool-calling ability. The term "agent" is used loosely in the industry for any AI assistant.

See: [Agentic AI](#); [Autonomy level](#); [Multi-agent system](#); [Tool calling \(function calling\)](#)

## Agent framework

Software infrastructure for building and deploying AI agents, often including tool integration, memory, and orchestration capabilities. Framework choice affects portability, vendor lock-in, and security posture; open source frameworks have different risk profiles than proprietary ones.

See: [Agentic AI](#); [Orchestration](#); [Tool calling \(function calling\)](#)

## [Agentic AI](#)

C R

An AI system that does one or both of the following: a) takes a larger request and breaks it down into smaller tasks for execution, and b) calls a tool (including another agent/system/model) and, based on the output, decides whether to provide a response, or continue (including changing the plan or calling another tool). Agentic AI systems require clear authorization boundaries, logging, human oversight mechanisms, and liability allocation for autonomous actions.

See: [Autonomy level](#); [Excessive agency](#); [Tool calling \(function calling\)](#)

## [Agents Rule of Two](#)

A security framework developed by Meta stating that AI agents should satisfy no more than two of the following three properties within a session: (A) processing untrustworthy inputs, (B) accessing sensitive systems or private data, and (C) changing state or communicating externally. Building on Simon Willison's Lethal Trifecta, the Rule of Two extends protection beyond data exfiltration to cover any state-changing action an agent might take, including examples like issuing refunds, modifying files, sending messages, or executing code. If a task requires all three properties, the agent should not operate autonomously and must include human-in-the-loop approval or equivalent supervision. The Rule of Two reflects the current consensus that prompt injection cannot be reliably detected or filtered, making architectural constraints the most practical defense for agentic AI systems.

See: [Agentic AI](#); [Excessive agency](#); [Human-in-the-loop](#); [Least privilege](#); [Lethal Trifecta](#); [Prompt injection](#)

## [Agent-to-Agent Protocol \(A2A\)](#)

C

An open source protocol enabling autonomous AI agents to discover, communicate, and collaborate with each other regardless of their underlying framework or vendor. A2A addresses the coordination layer: how agents delegate tasks, share capabilities, and track outcomes. Note: The A2A effort merged with IBM's Agent Communication Protocol (ACP).

See: [Agentic AI](#); [Model Context Protocol](#); [Multi-agent system](#); [Orchestration](#)

## [AI Bill of Materials \(AIBOM\)](#)

## C

An inventory documenting the components, dependencies, and provenance of an AI system, analogous to an SBOM for software. An AIBOM typically covers the model(s) in an AI system and their versions, fine-tuning datasets, adapters, system prompts, guardrails, embedding models, vector databases, tools and plugins, and third-party services—the full CORE of the deployed system. AIBOMs support license compliance, supply chain security, incident response, and regulatory documentation requirements. They are increasingly requested in enterprise procurement and expected for high-risk AI systems.

See: [AI system](#); [CORE](#); [License compliance](#); [Model card](#); [Supply chain security](#)

## [AI governance](#)

### C R

The policies, processes, roles, and controls for managing AI risk and ensuring responsible development and deployment. Regulators, customers, and boards increasingly expect AI governance programs, and support reasonable care arguments and compliance demonstrations.

See: [AI policy](#); [Model registry](#); [NIST AI RMF \(AI Risk Management Framework\)](#); [Risk assessment](#)

## [AI model](#)

See: [Model](#)

## [AI policy](#)

### R

Internal organizational policy or external governmental policy addressing development, deployment, procurement, and use of AI systems. Policies commonly cover permitted uses, data handling, human oversight, documentation, and incident response.

See: [Acceptable Use Policy](#); [AI governance](#); [Monitoring](#); [Risk assessment](#)

## [AI safety](#)

Research and practices aimed at reducing harmful or unintended behavior of AI systems. In enterprise, procurement, and policy contexts, “AI safety” may refer to model training and post-training methods (e.g., RLHF), system-level guardrails, abuse monitoring, and evaluation programs tied to a stated threat model.

See: [Alignment](#); [Guardrails](#); [Red teaming](#); [Safety policy](#)

## [AI system](#)

### C R

The complete deployed software system, including all dataflow-affecting components, models, guardrails, control systems, accessible resources, allowed operations, and interfaces—not just the underlying model(s). An AI system's risk profile depends not only on model capabilities but on its CORE: the Components it comprises, the Operations it can perform, the Resources it can reach, and the Execution dataflow connecting them. Regulatory obligations under the EU AI Act and sector rules typically attach to the "AI system" as deployed, making it important to distinguish from "model" in contracts and governance.

See: [CORE](#); [Deployer](#); [High-risk AI system](#); [Model](#); [Provider](#)

### [Algorithmic discrimination](#)

L R

Differential treatment or outcomes produced by an AI system that may be associated with protected characteristics or other legally relevant categories. Depending on jurisdiction and use case, this concept can be relevant to civil rights, consumer protection, employment, housing, credit, and sector-specific obligations. Documentation and testing may be used to evaluate risk and compliance.

See: [Bias](#); [Disparate impact](#); [Fairness](#); [High-risk AI system](#)

### [Algorithmic Disgorgement \(Model Deletion\)](#)

R

A regulatory remedy requiring deletion of unlawfully obtained data and, in some cases, deletion of models or derived artifacts trained using that data (including weights, checkpoints, or fine-tunes). It may be sought where deleting the underlying dataset alone would not address downstream effects.

See: [Machine Unlearning](#); [Prohibited AI practices](#); [Training data](#)

### [Alignment](#)

R

Ensuring AI system objectives and behaviors match human values, intentions, and instructions. Alignment encompasses both technical measures and governance, and affects safety representations and liability exposure.

See: [Constitutional AI](#); [Safety policy](#); [Value alignment](#)

### [Annotation](#)

The process of labeling data for training or evaluation; also called "tagging" or "labeling." Annotation involves human labor (often outsourced), content exposure, and quality control issues; IP questions arise for labeled datasets.

See: [Ground truth](#); [Labeling](#); [Training data](#)

## Anonymization

R

A form of data processing intended to make data no longer identifiable to any individual, typically requiring that re-identification is not reasonably likely given available means. Legal definitions and thresholds vary by jurisdiction and context.

See: [De-identification](#); [Personal data](#); [Pseudonymization](#)

## Answer relevance

A retrieval-augmented generation (RAG) evaluation measure assessing how well an answer addresses the user's question, often compared to a reference or judged by humans/LLMs. It is frequently used alongside context relevance and groundedness.

See: [Context relevance](#); [Evaluation \(evals\)](#); [Groundedness](#)

## Application Programming Interface (API)

C

A structured interface allowing software to communicate with other software programmatically, as distinguished from a human-facing interface like a website or app. In the AI context, API access means integrating AI capabilities directly into applications, workflows, or products through code rather than through a chat interface. API terms typically differ significantly from consumer terms: they often permit broader commercial use and integration but impose rate limits, usage-based pricing, data handling obligations, and restrictions on downstream redistribution. Key contractual issues include whether outputs can be used to train competing models, what data is logged and retained, SLA commitments, and how usage is metered and billed.

See: [Endpoint](#); [Rate limiting](#); [Service Level Agreement / Service Level Objective](#); [Usage data / telemetry](#); [Zero Data Retention](#)

## Approval workflow

C

A process requiring human authorization before certain actions or outputs are executed or released. Approval workflows are key controls for high-risk agentic AI systems.

See: [Autonomy level](#); [Change control](#); [Human-in-the-loop](#)

## Architecture

The high-level structure of a model or AI system and how components interact (e.g., model + retrieval + tools + guardrails + monitoring). Architecture choices influence performance, security, privacy, and auditability. One way to identify the parts of an architecture is by using the mnemonic CORE - Components, Operations, Resources, and Execution.

See: [AI system](#); [CORE](#); [Guardrails](#); [Tool calling \(function calling\)](#)

## [Artificial General Intelligence \(AGI\)](#)

A non-standard term used to describe hypothetical AI with broad, human-level capability across many domains. The term is used inconsistently in technical and marketing contexts; many current systems are better described as foundation models or general-purpose AI rather than “general intelligence.” In diligence and policy discussions, “AGI” sometimes signals discussion of frontier capability thresholds and risk controls.

See: [Capabilities](#); [Foundation model](#); [Narrow AI](#)

## [Attention mechanism](#)

T

The method by which transformer models determine which parts of an input matter for producing each part of an output. When generating the next word, the model assigns weights to every previous token, attending more to relevant context and less to irrelevant text. Attention is not comprehension; the model is computing statistical relevance, not understanding meaning. The attention patterns can sometimes be examined to understand why a model produced particular outputs, though this interpretability has limits.

See: [Context window](#); [Self-attention](#); [Transformer](#)

## [Audit](#)

C R

Independent examination of AI systems, processes, or controls for compliance, performance, or risk assessment. Audit rights are common in enterprise contracts. Third-party AI audits are emerging as a governance and compliance tool.

See: [AI governance](#); [Conformity assessment](#); [Evaluation \(evals\)](#); [SOC 2](#)

## [Audit log](#)

A record of system events that supports tracing actions, changes, and access (e.g., who accessed data, what tools were called, what model version ran). Audit logs are commonly used for security investigations, compliance, and dispute resolution.

See: [Access control](#); [Change control](#); [Logging](#); [Monitoring](#)

## [AUP](#)

See: [Acceptable Use Policy](#)

## [Authentication](#)



Verification of user or system identity before granting access to AI services or data. Authentication controls are baseline security requirements; failures can create breach liability and confidentiality exposure.

See: [Access control](#); [Security](#)

## Authorization

The process of determining what an authenticated user, service, or agent is allowed to do (permissions), often implemented via roles, policies, or scopes. Authorization differs from authentication (identity verification).

See: [Access control](#); [Authentication](#); [Least privilege](#)

## Autoencoder

T

A neural network that learns compressed representations by encoding inputs and then reconstructing them. Used in anomaly detection, data compression, and generative models; relevant when understanding technical architecture of certain AI systems.

See: [Decoder](#); [Encoder](#); [Latent space](#)

## Automated decision-making (ADM)

R

Decisions made by AI systems with limited or no human involvement, often subject to regulatory requirements. GDPR Article 22, various US state laws, and sector regulations impose requirements on ADM including rights to explanation and human review.

See: [Explainability](#); [Human-in-the-loop](#); [Right to explanation](#)

## Autonomy level

C R

The degree of independence an AI system has in making and executing decisions without human oversight. Autonomy level is central to risk assessment, liability allocation, and regulatory compliance for agentic systems.

See: [Agentic AI](#); [Approval workflow](#); [Human-in-the-loop](#)

## Availability

C

The percentage of time a system is operational and accessible, often defined in SLAs. Availability commitments for AI-dependent systems often account for model-specific issues (capacity, rate limits, provider outages).

See: [Business continuity](#); [Service Level Agreement / Service Level Objective](#); [Uptime](#)

## B

### Backdoor (model backdoor / trojan)

A hidden behavior in a model or system that activates under specific triggers (e.g., particular phrases, patterns, or inputs), causing unintended outputs or actions. Backdoors are discussed in model security, supply chain risk, and red teaming contexts.

See: [Adversarial attack](#); [Data poisoning](#); [Model supply chain](#); [Red teaming](#)

### Backpropagation

The algorithm used to train neural networks by computing how much each weight contributed to prediction errors, then adjusting weights to reduce those errors. Backpropagation is how models "learn" from training data; errors propagate backward through the network, and weights are updated accordingly.

See: [Gradient descent](#); [Training](#); [Weights](#)

### Base model

T

A pre-trained model before any fine-tuning or customization; the starting point for downstream adaptations.

See: [Fine-tuning](#); [Foundation model](#); [Pre-training](#)

### Batch size

The number of training examples processed together before updating model weights; a hyperparameter affecting training dynamics. Affects training resource requirements and costs.

See: [Epoch](#); [Hyperparameter](#); [Training](#)

### Benchmark

C T

A standardized test for measuring model performance on specific tasks, enabling comparison across models. Benchmark claims in marketing often specify which benchmark, version, and conditions; benchmarks may not reflect real-world performance.

See: [Accuracy](#); [Evaluation \(evals\)](#)

### Bias

L R

Systematic differences in model behavior or error rates that correlate with particular features, groups, or contexts. Bias is not limited to interactions with humans; it reflects

the degree to which the distribution of features in training data matches the distribution in production. A model trained primarily on certain populations, document types, or scenarios will perform differently on others. In regulated contexts, bias testing and mitigation records are used to assess compliance posture.

See: [Algorithmic discrimination](#); [Disparate impact](#); [Fairness](#)

## [Biometric data](#)

### R

Data derived from physical or behavioral characteristics used for identification (face, voice, fingerprint, gait). Biometric data triggers heightened obligations under BIPA, GDPR, state privacy laws, and the EU AI Act; AI systems processing biometrics require special controls.

See: [Multimodal model](#); [Personal data](#); [Privacy](#)

## [Black box](#)

A system whose internal decision process is difficult to interpret or explain in a human-understandable way. The term is used in technical, governance, and legal contexts when evaluating transparency, accountability, and auditability.

See: [Explainability](#); [Model card](#); [System card](#); [XAI](#)

## [Business continuity](#)

### C

Plans and capabilities for maintaining operations during disruptions, including AI system failures or provider outages. AI-dependent workflows need continuity planning; contracts commonly address provider failures, model deprecation, and data portability.

See: [Availability](#); [Portability](#); [Vendor lock-in](#)

## **C**

## [Capabilities](#)

The tasks and performance characteristics a model or system can reliably support (e.g., summarization, coding, extraction, tool use). “Capabilities” is often tied to evaluation results and can be discussed in procurement, marketing, and policy settings.

See: [Benchmark](#); [Evaluation \(evals\)](#); [Intended use](#); [Model capability](#)

## [Chain of Thought \(CoT\)](#)

### T

A prompting technique that encourages models to show intermediate reasoning steps, often improving accuracy on complex tasks. CoT reasoning may provide some

transparency into model "thinking" but is not a substitute for true explainability; reasoning traces may be fabricated.

See: [Explainability](#); [Prompting](#); [Reasoning model](#)

## [Change control](#)

C

Processes governing modifications to AI systems, including model updates, prompt changes, and configuration adjustments. Change control is essential for regulated deployments; contracts commonly specify notice, approval, testing, and rollback requirements.

See: [Model drift](#); [Model update](#); [Version pinning](#)

## [Chat model](#)

T

An LLM configured or post-trained for conversational interaction (e.g., instruction following, dialogue safety behaviors). Chat models are often accessed through chat-completion interfaces and may differ from base models in behavior and safety characteristics. ChatGPT was the first widely known chat model and remains the best-known example.

See: [Instruction tuning](#); [Large Language Model](#); [System prompt](#)

## [Checkpoint](#)

A saved state of model weights during training or fine-tuning, used to resume training or to preserve intermediate versions. In training agreements, checkpoints may be deliverables; agreements often specify ownership, retention, access controls, and permitted reuse.

See: [Model artifact](#); [Training](#); [Weights](#)

## [Chemical, Biological, Radiological, Nuclear \(CBRN\)](#)

Categories of weapons or hazardous materials; AI systems' potential to assist with CBRN threats is a key safety concern. CBRN-related content is typically prohibited in AUPs; frontier AI safety evaluations specifically test for CBRN assistance capabilities.

See: [Dual-use](#); [Prohibited AI practices](#); [Safety evaluation](#)

## [Child Sexual Abuse Material \(CSAM\)](#)

R

Illegal content depicting sexual abuse of minors. Many AI services implement controls to detect, block, and report CSAM due to legal obligations and platform policy

requirements; specific duties, reporting pathways, and retention requirements vary by jurisdiction and service design.

See: [Content filtering](#); [Prohibited AI practices](#); [Safety policy](#)

## Chunking

Dividing documents into smaller segments for processing within context window limits or for retrieval purposes. Chunking strategies affect retrieval accuracy and completeness; relevant when assessing whether AI systems properly considered full documents.

See: [Context window](#); [Truncation](#)

## Citation

A reference to a source used to support a statement or output (e.g., a retrieved document chunk in RAG, or a legal citation to authority). In AI systems, “citations” may be generated automatically and can be incorrect or incomplete unless the system is designed to capture provenance.

See: [Grounding](#); [Hallucination](#); [Source attribution](#)

## Classification

The task of assigning inputs to predefined categories (e.g., spam detection, sentiment analysis, content moderation). Classification errors have different consequences depending on the application; false positives and negatives have different risk profiles.

See: [Accuracy](#); [Discriminative model](#)

## Cloud provider

### C

A vendor providing cloud infrastructure and managed services (compute, storage, networking) used to train or operate AI systems. Cloud provider terms commonly affect data residency, security controls, incident response, and subprocessors.

See: [Data residency](#); [Hosting](#); [Security addendum](#); [Subprocessor](#)

## Clustering

Grouping similar items together based on their features without predefined labels. Clustering can produce de facto sensitive inferences (grouping by health, demographics) even without explicit attributes.

See: [Embedding](#); [Privacy](#); [Unsupervised learning](#)

## CNN

See: [Convolutional Neural Network](#)

## [Coalition for Content Provenance and Authenticity \(C2PA\)](#)

A technical standard for embedding provenance information in digital content, including AI-generated content. C2PA and similar standards support content authenticity verification; increasingly relevant for deepfake detection and evidence authentication.

See: [Content provenance](#); [Deepfake](#); [Watermarking](#)

## [Components](#)

In the CORE framework, the functional elements that comprise an AI system: models, adapters, guardrails, databases, APIs, connectors, human review steps, and other nodes through which data flows. Each component has properties relevant to governance: its provider or origin, the operations it performs, the resources it accesses, and how it transforms the data flowing through it. A list of components is the minimum information needed for an AIBOM.

See: [Adapter](#); [AI system](#); [CORE](#); [Guardrails](#); [Model](#)

## [Compute](#)

C R

The computational resources (processing power, memory, storage) required to train and run AI systems; also a regulatory concept. Compute thresholds trigger reporting requirements under various Executive Orders, state laws, and the EU AI Act. Compute access is a key factor in AI capabilities.

See: [Export controls](#); [Inference](#); [Training](#)

## [Computer vision](#)

T

A field of AI focused on interpreting and generating information from images and video (e.g., object detection, segmentation, captioning). Many modern systems use multimodal models that combine vision and language capabilities.

See: [CNN](#); [Multimodal](#); [Vision-language model](#)

## [Confidential computing](#)

Hardware-supported techniques protecting data while in use (in memory) using secure enclaves. Confidential computing may support stronger security representations for processing sensitive data in cloud environments.

See: [Encryption](#); [Security](#); [Trusted execution environment](#)

## [Confidential information](#)

## C

Information protected under contract or law from unauthorized use or disclosure; in AI contexts often includes prompts, outputs, logs, and fine-tuning artifacts.

See: [Logging](#); [Trade secret](#); [Usage data / telemetry](#)

### Confidentiality

Protection of non-public information from unauthorized disclosure. In AI contracting, confidentiality terms often address prompts, outputs, logs, and resources available to the AI system such as connected enterprise data.

See: [Access control](#); [Confidential information](#); [Logging](#)

### Conformity assessment

## R

A structured assessment process (often regulatory) to verify that a system meets specified requirements before deployment. EU AI Act and some sector regulations require conformity assessments for higher-risk AI; determine who performs them and what evidence is retained.

See: [EU AI Act](#); [High-risk AI system](#); [Technical documentation](#)

### Connector

An integration that pulls content from enterprise systems (e.g., SharePoint, Google Drive, Slack) into an AI system for retrieval or context. Connectors expand the data-access surface area; permissions, logging, and retention practices are commonly evaluated to reduce privilege, confidentiality, and privacy risk.

See: [Access control](#); [CORE](#); [Knowledge base](#); [Operations](#); [Resources](#)

### Consent

## L R

Agreement by an individual to a specified use of their data, likeness, or content. In AI contexts, consent may be relevant to data collection, biometric processing, voice cloning, and digital replicas; requirements vary by jurisdiction and sector.

See: [Biometric data](#); [Digital Replica](#); [Personal data](#); [Right of Publicity](#)

### Constitutional AI

An alignment approach using a set of principles ("constitution") to guide model behavior, often using AI-generated feedback. Constitutional policies can be relevant to content and safety representations; request documentation for high-stakes use cases.

See: [Alignment](#); [Safety policy](#)

## Content filtering

C

Automated detection and blocking or transformation of disallowed content. Filtering affects safety claims, regulatory compliance, intellectual property infringement, and AUP enforcement; raises false positive/negative issues.

See: [Guardrails](#); [Moderation](#); [Safety policy](#)

## Content provenance

R

Information describing the origin and transformation history of content, including whether AI was involved in creation. Provenance supports authenticity, IP compliance, and consumer transparency; increasingly relevant for evidence authentication and misinformation disputes.

See: [Deepfake](#); [Metadata](#); [Watermarking](#)

## Context overload attack

An adversarial technique overloading the prompt with excessive tokens to predispose models to a vulnerable state. A specific prompt injection variant; relevant to security testing and incident analysis.

See: [Adversarial attack](#); [Context window](#); [Prompt injection](#)

## Context relevance

A RAG quality metric measuring whether retrieved context contains information pertinent to the user's query. Poor context relevance can cause unreliable outputs; relevant when evaluating RAG system performance claims.

See: [Answer relevance](#); [Groundedness](#)

## Context window

C T

The maximum number of tokens a model can consider at once, encompassing both the input and the output being generated. Context windows have expanded dramatically—from about 4,000 tokens in early GPT-4 to over 1 million tokens in some current models—but limits still matter. When input exceeds the context window, content is truncated, often without notification to the user. Context window size is distinct from how well a model uses that context; performance often degrades on information buried in the middle or nearer to the end of long inputs.

See: [Token](#); [Truncation](#)

## Continual learning



A learning paradigm where AI systems incrementally learn from new data while preserving prior knowledge (avoiding "catastrophic forgetting"). Continual learning systems may evolve in ways that affect prior representations about behavior; governance often addresses ongoing changes.

See: [Model drift](#); [Model update](#); [Training](#)

## [Continuous learning](#)

See: [Continual learning](#)

## [Contract](#)

C

A legally binding agreement defining rights, obligations, and remedies between parties. In AI contexts, contracts often define scope of use, data rights, confidentiality, IP terms, warranties, indemnities, SLAs, and change control.

See: [Indemnity](#); [Limitation of liability](#); [SLA/SLO](#); [Terms of Service](#)

## [Controller / processor](#)

C R

Privacy law roles: controller determines purposes/means of processing; processor processes on behalf of controller. AI vendors often characterize themselves as processors; customers may require controls on model training and subprocessing consistent with that role.

See: [DPA](#); [Personal data](#); [Subprocessor](#)

## [Controls](#)

Technical or organizational measures used to achieve defined objectives (e.g., security controls, privacy controls, safety controls). In governance and audits, "controls" are often documented, tested, and monitored. In contrast to regular GRC systems, AI controls need to be implemented with technical measures, frequently as components external to the model.

See: [Audit](#); [Monitoring](#); [Privacy-enhancing technology](#); [Security](#)

## [Convolutional Neural Network \(CNN\)](#)

T

A neural network architecture designed for processing grid-like data (images), using convolutional layers to extract features. CNNs power most computer vision applications including facial recognition and autonomous vehicles.

See: [Computer vision](#); [Deep learning](#); [Neural network](#)

## [Copyleft license](#)

## C

Open source licenses that may require distributing source code or licensing downstream when distribution triggers occur (e.g., GPL). Copyleft obligations can create compliance risk when AI systems distribute software or embed licensed components.

See: [License compatibility](#); [Open source](#)

## Copyright

### L

A body of law protecting original works of authorship fixed in a tangible medium, granting exclusive rights (e.g., reproduction, distribution, derivative works) subject to limitations and exceptions. In AI discussions, copyright commonly arises with training data provenance, output ownership, and infringement/fair use analysis.

See: [Copyright infringement \(AI context\)](#); [Fair use](#); [Output](#); [Training data](#)

## Copyright infringement (AI context)

### L

Claims that AI outputs reproduce copyrighted material or that AI training constituted infringement. Central issue in pending litigation; training data provenance, output filtering, and indemnification are key contractual topics.

See: [Fair use](#); [IP indemnity](#); [Substantial similarity](#); [Training data](#)

## CORE

### C

A framework and mnemonic device for analyzing and documenting AI systems by mapping their Components, Operations, Resources, and Execution dataflow. CORE represents AI systems as directed graphs where data flows through connected elements, enabling policy evaluation, compliance tracking, and risk assessment over an entire AI system.

See: [AI governance](#); [AI system](#); [Components](#); [Execution](#); [Operations](#); [Resources](#)

## Cosine similarity

A metric comparing embeddings by measuring the angle between vectors; higher similarity implies closer semantic meaning. Relevant when understanding RAG retrieval mechanics; similarity thresholds affect retrieval defensibility.

See: [Embedding](#); [Semantic search](#); [Vector database](#)

## CoT

### T

See: [Chain of Thought \(CoT\)](#)

## [Cross-border data transfer](#)

C R

Moving personal data across national borders, triggering transfer mechanisms and restrictions. AI vendors may route prompts/logs across regions; DPAs and data residency terms often match actual architecture.

See: [Data residency](#); [DPA](#); [Privacy](#)

## [CVE](#)

C

See: [Vulnerability / CVE](#)

## [Cyber Resilience Act \(CRA\)](#)

R

EU Regulation 2024/2847 establishing mandatory cybersecurity requirements for "products with digital elements" (hardware and software connected to devices or networks) sold in the EU market. The CRA entered into force December 2024, with full applicability by December 2027. It requires manufacturers to ensure products are secure by design, maintain vulnerability management throughout the product lifecycle, provide security updates, and report actively exploited vulnerabilities. Products are classified by risk level (critical, important, or default), with higher-risk products requiring third-party conformity assessment. The CRA applies to most software including AI systems and their components; it intersects with the EU AI Act (which addresses AI-specific risks) and requires SBOM-like documentation of components. Open source software developed outside commercial activity is generally exempt, though commercial products incorporating open source remain in scope.

See: [EU AI Act](#); [Security](#); [Supply chain security](#); [Vulnerability / CVE](#)

# **D**

## [Data annotation](#)

See: [Annotation](#)

## [Data augmentation](#)

Techniques that expand or vary training data to improve model generalization (e.g., transformations, paraphrases, synthetic examples). Augmentation can affect performance and bias characteristics depending on how it is applied.

See: [Bias](#); [Synthetic data](#); [Training data](#)

## Datacenter

A facility housing computing infrastructure (servers, storage, networking) used for training and operating AI systems. Datacenter location and subcontracting can matter for data residency, security, and resilience.

See: [Business continuity](#); [Cloud provider](#); [Data residency](#)

## Data classification

C R

An organizational scheme for labeling resource data by sensitivity and handling requirements (e.g., public, internal, confidential, regulated). Data classification commonly drives AI policy decisions about prompts, connectors, logging, and vendor use.

See: [Confidential information](#); [Connector](#); [Data governance](#); [Logging](#); [Resources](#)

## Data controller

See: [Controller / processor](#)

## Data drift

Changes over time in the distribution or characteristics of input data (or user behavior) that can degrade model performance. Data drift is often monitored alongside model drift (changes in the model itself).

See: [Evaluation \(evals\)](#); [Model drift](#); [Monitoring](#)

## Data governance

C R

Organizational processes and controls for managing data quality, access, lineage, retention, and compliance. Data governance often sets the baseline for AI policy, connector use, logging, and training restrictions.

See: [Access control](#); [Data classification](#); [Retention](#); [Training data](#)

## Data leakage

C

Unintended disclosure of training data, prompts, or other sensitive information through model outputs. Data leakage can breach confidentiality, privacy, and security obligations; memorization and extraction attacks are key concerns.

See: [Membership inference](#); [Memorization](#); [Model inversion](#)

## Data Loss Prevention (DLP)

Technologies and policies preventing unauthorized data transmission or exposure. DLP tools can help prevent sensitive data from being sent to AI systems; relevant for shadow AI governance.

See: [Confidential information](#); [Security](#); [Shadow AI](#)

## [Data minimization](#)

R

A privacy principle requiring collection and retention of only data necessary for specified purposes. Data minimization applies to AI training, inference logging, and improvement uses; conflicts with desires for comprehensive data may arise.

See: [Privacy by design](#); [Purpose limitation](#)

## [Data poisoning](#)

An adversarial attack inserting or modifying training data to compromise model behavior. Data poisoning attacks can undermine model integrity; supply chain security and data provenance controls are relevant defenses.

See: [Adversarial attack](#); [Supply chain security](#); [Training data](#)

## [Data privacy attack](#)

Attacks designed to gain access to sensitive information in training data, including membership inference and model inversion. Privacy attacks demonstrate that training data can sometimes be extracted; relevant to privacy representations and security controls.

See: [Data leakage](#); [Membership inference](#); [Model inversion](#)

## [Data Processing Addendum \(DPA\)](#)

C

A contract addendum addressing personal data processing requirements, typically required by privacy regulations. DPAs often address AI-specific issues: training uses, inference logging, subprocessors, model improvement, and cross-border transfers.

See: [Controller / processor](#); [Personal data](#); [Subprocessor](#)

## [Data processor](#)

See: [Controller / processor](#)

## [Data Protection Impact Assessment \(DPIA\)](#)

R C

An assessment required under GDPR Article 35 before processing likely to result in high risk to individuals' rights and freedoms, including systematic profiling, large-scale

processing of sensitive data, and systematic monitoring. DPIAs must describe the processing, assess necessity and proportionality, identify risks, and specify mitigations. AI systems frequently trigger DPIA requirements due to automated decision-making, profiling, and processing at scale. Unlike FRIAs (which address broader fundamental rights), DPIAs focus specifically on data protection risks.

See: [Automated decision-making](#); [Personal data](#); [Privacy](#)

### [Data Protection Officer \(role\) \(DPO\)](#)

R

A designated individual responsible for overseeing GDPR compliance, advising on data protection obligations, and serving as the point of contact for supervisory authorities and data subjects. Not to be confused with Direct Preference Optimization (DPO) (training).

See: [Controller / processor](#); [DPO](#); [Privacy](#)

### [Data provenance](#)

Documentation of data origin, collection methods, transformations, and chain of custody. Provenance is essential for IP compliance, bias assessment, and regulatory documentation; request provenance information for training data.

See: [Content provenance](#); [Dataset documentation](#); [Training data](#)

### [Data reconstruction](#)

A privacy attack designed to reconstruct sensitive information from training data through model queries. Reconstruction attacks demonstrate privacy risks from training; relevant to data sensitivity assessments.

See: [Data privacy attack](#); [Memorization](#); [Model inversion](#)

### [Data residency](#)

C R

Requirements that data be stored or processed in specific geographic locations. Data residency requirements may constrain AI architecture choices; verify that vendor infrastructure meets residency requirements.

See: [Cross-border data transfer](#); [DPA](#); [Subprocessor](#)

### [Data retention](#)

C

Policies governing how long data is kept before deletion. Retention of prompts, outputs, and logs creates ongoing risk; align retention with purpose limitation and deletion rights.

See: [Logging](#); [Privacy](#); [Records retention](#)

## Data security

C R

Safeguards protecting data confidentiality, integrity, and availability, including access controls, encryption, logging, incident response, and secure development practices. In AI systems, data security applies to prompts, outputs, logs, embeddings, and connected enterprise data.

See: [Access control](#); [Encryption](#); [Security](#); [Security addendum](#)

## Data segregation

C

Separation of customer data or environments to prevent unauthorized access or cross-tenant leakage (logical or physical). Segregation is relevant to multi-tenant AI services, RAG indexes, and logging systems.

See: [Access control](#); [Confidentiality](#); [Data residency](#); [Multi-tenant](#)

## Dataset documentation

T

Structured information about a dataset's contents, collection, limitations, and intended uses (e.g., datasheets, data cards). Dataset documentation supports IP diligence, bias assessment, and regulatory compliance; request it for training datasets.

See: [Data provenance](#); [Model card](#); [Training data](#)

## Data subject

R

An identified or identifiable individual whose personal data is processed. In privacy frameworks (e.g., GDPR and state privacy laws), data subjects may have rights (access, deletion, correction, objection, portability, etc.), and organizations typically implement processes to respond to those rights.

See: [Personal data](#); [Privacy](#)

## Decisioning

C

Using AI outputs to make or inform consequential decisions about individuals or transactions. Decisioning triggers regulatory obligations, liability exposure, and fairness requirements; distinguish advisory outputs from decisioning.

See: [Adverse action](#); [Automated decision-making](#); [High-risk AI system](#)

## Decoder

## T

A neural network component that reconstructs outputs from compressed representations. Decoders are the generative component in many AI architectures.

See: [Autoencoder](#); [Encoder](#); [Transformer](#)

### Decoder-only model

## T

A transformer model using only the decoder, suited for text generation tasks (e.g., GPT family). Most modern LLMs are decoder-only architectures.

See: [Encoder-only model](#); [LLM](#); [Transformer](#)

### Deepfake

## L R

Synthetic media created using AI to depict people saying or doing things they did not actually say or do. Deepfakes raise defamation, fraud, election, NCII, and evidence authentication issues; detection and provenance tools are evolving.

See: [Content provenance](#); [Generative AI \(GenAI\)](#); [NCII](#); [Synthetic media](#)

### Deep learning

Machine learning using neural networks with multiple layers, enabling detection of complex patterns in data. The term is often used interchangeably with "AI" in business contexts, though technically it refers to a specific architectural approach. The "deep" in deep learning refers to the number of layers between input and output, not to any quality of understanding.

See: [Machine learning](#); [Neural network](#); [Training](#)

### De-identification

## R

Techniques that remove or obscure identifiers to reduce the ability to link data to a specific individual. De-identification is commonly evaluated based on the risk of re-identification given available auxiliary data, threat models, and technical safeguards; requirements and standards vary by law and context.

See: [Anonymization](#); [Personal data](#); [Pseudonymization](#)

### Deletion

Removal of data from systems and backups according to defined policies and technical processes. In privacy and AI contexts, deletion can relate to data subject rights, retention policies, and requests to remove training data or derived artifacts.



See: [Algorithmic Disgorgement \(Model Deletion\)](#); [Data subject](#); [Machine Unlearning](#); [Retention](#)

## [Deliverables](#)

C

Items a party is obligated to provide under an agreement (e.g., fine-tuned model, adapter, documentation, evaluation results, training logs, or a deployed service). In AI projects, deliverables are often defined to clarify ownership, acceptance, and maintenance responsibilities.

See: [Acceptance criteria](#); [Documentation](#); [Model artifact](#)

## [Deployer](#)

R

A party that deploys an AI system for use, as distinguished from the provider/developer. EU AI Act and other frameworks allocate different obligations to providers vs. deployers; determine your role and resulting duties.

See: [AI system](#); [EU AI Act](#); [Provider](#)

## [Deterministic](#)

A system property where the same input always produces the same output. Traditional software is deterministic by design. AI systems can in theory be configured for deterministic behavior, though hardware and infrastructure variations may still introduce variability. Deterministic operation supports reproducibility, testing, audit, and regulatory compliance, but may reduce output quality or diversity compared to default settings.

See: [Non-deterministic](#); [Reproducibility](#); [Sampling](#); [Temperature](#)

## [Developer prompt](#)

T

Instructions provided by an application or developer that guide model behavior for a specific product or workflow (e.g., formatting, tool use rules, safety constraints). Developer prompts typically have lower priority than system prompts but higher priority than user prompts.

See: [Instruction hierarchy \(prompt precedence\)](#); [System prompt](#); [User prompt](#)

## [Development practices](#)

The processes and standards used to build and maintain software and AI systems (e.g., secure coding, testing, review, documentation, incident response). These practices are often described in security questionnaires, audits, or contractual representations.

See: [Documentation](#); [Secure development lifecycle](#); [Testing](#)

## DevOps

Practices combining software development and operations to improve deployment frequency, reliability, and monitoring (e.g., CI/CD, infrastructure as code). In AI settings, DevOps is often paired with MLOps for model lifecycle management.

See: [Change control](#); [MLOps](#); [Monitoring](#)

## Differential privacy

R

A privacy technique that adds controlled statistical noise to limit what can be inferred about any individual record from an aggregate output. Differential privacy is discussed in privacy engineering, dataset release, and sometimes in training/evaluation contexts.

See: [De-identification](#); [Personal data](#); [Privacy-enhancing technology](#)

## Diffusion model

T

A generative AI architecture that creates outputs by iteratively denoising random noise, commonly used for image generation. Diffusion models power DALL-E, Midjourney, Stable Diffusion; they raise distinctive copyright and provenance questions for images.

See: [Generative AI \(GenAI\)](#); [Image generation](#); [Text-to-image](#)

## Digital Replica

L R

A term used in right-of-publicity and synthetic media discussions to describe a computer-generated representation of a person's image, voice, or likeness. Applicable requirements and remedies vary by jurisdiction and may depend on consent, context, and whether the replica is used for commercial or deceptive purposes.

See: [Deepfake](#); [Right of Publicity](#); [Synthetic media](#)

## Direct Preference Optimization (training) (DPO)

T

A training technique that aligns models using preference data without requiring a separate reward model. DPO is an alternative to RLHF for alignment training. Not to be confused with Data Protection Officer (DPO) (role).

See: [Alignment](#); [Training](#)

## Disclosure

Communication of information to a user, counterparty, regulator, or the public (e.g., about AI use, limitations, data practices, or incidents). Disclosure duties can arise from contracts, consumer protection rules, sector regulations, or internal governance.

See: [Documentation](#); [Notice](#); [Transparency](#)

### [Discriminative model](#)

A model that classifies inputs or distinguishes between categories rather than generating new content. Discriminative models (classifiers, detectors) have different risk profiles than generative models; errors are often binary.

See: [Classification](#); [Generative AI \(GenAI\)](#)

### [Disparate impact](#)

L R

Facially neutral practices that disproportionately affect protected groups. Disparate impact analysis applies to AI systems used in employment, credit, housing; testing often assess group-level outcomes.

See: [Algorithmic discrimination](#); [Bias](#); [Fairness](#)

### [Distillation](#)

T

Training a smaller "student" model to mimic a larger "teacher" model's behavior by learning from the teacher's outputs rather than the original training data. Distillation does not require access to the teacher model's weights, only the ability to query it and observe its outputs. This creates trade secret and competitive concerns: even without sharing weights, unguarded API access may allow third parties to replicate proprietary model capabilities. Distillation can also transfer copyrighted expression if the teacher's outputs are used as training data. Many model licenses and API terms of service explicitly prohibit using outputs to train other models.

See: [AUP](#); [Knowledge transfer](#); [Model compression](#); [Model extraction](#); [Trade secret](#); [Training](#); [Weights](#)

### [Distributional robustness](#)

A model's ability to perform equitably across the range of possible inputs, including rare or "long-tail" cases. Distributional robustness affects fairness and reliability; models may fail on underrepresented populations or edge cases.

See: [Evaluation \(evals\)](#); [Fairness](#); [Robustness](#)

### [Documentation](#)

Written materials describing a model, system, dataset, or process (e.g., model cards, system cards, technical specs, policies, runbooks). Documentation is commonly used for governance, audits, contracting, and incident response.

See: [Audit](#); [Dataset documentation](#); [Model card](#); [System card](#)

## DPA

See: [Data Processing Addendum](#)

## DPO

### T

See: [Data Protection Officer \(role\)](#); [Direct Preference Optimization \(training\)](#)

## Dual-use

### R

Technology that has both legitimate and potentially harmful applications. Dual-use considerations affect export controls, safety evaluations, and responsible deployment decisions for frontier AI.

See: [Dual-use foundation model](#); [Export controls](#)

## Dual-use foundation model

### R T

A foundation model that exhibits or could be modified to exhibit high levels of performance at tasks posing serious security risks. Definition from Executive Order 14110; triggers reporting requirements and enhanced safety obligations.

See: [Executive Orders on AI](#); [Foundation model](#)

## E

## Edge deployment

Running AI models on local devices rather than cloud servers. Edge deployment affects data residency, latency, security, and update mechanisms; may be preferred for privacy-sensitive applications.

See: [Data residency](#); [Latency](#); [On-prem deployment](#); [Small Language Model](#)

## E-discovery

### L

The process of identifying, preserving, collecting, and producing electronically stored information (ESI) in litigation or investigations. AI system logs, prompts, outputs, tool calls, and model/version records can be relevant ESI depending on the matter.

See: [Audit log](#); [Litigation hold](#); [Logging](#); [Retention](#)

## [Embedding](#)

C T

A list of numbers (vector) representing the meaning of text, images, or other content in a form that enables mathematical comparison. Two pieces of text with similar meanings will have similar embeddings, allowing systems to find semantically related content even without shared keywords. Embeddings power RAG retrieval: when a user asks a question, the system converts it to an embedding and finds stored documents with nearby embeddings.

See: [Semantic search](#); [Vector](#); [Vector database](#)

## [Embedding model](#)

T

A model specifically designed to generate embeddings for retrieval and similarity tasks. Embedding model selection affects retrieval quality and privacy risk (what information is encoded).

See: [Embedding](#); [Vector database](#)

## [Emergent capabilities](#)

Unexpected abilities that appear in larger models without explicit training for those tasks. Emergent capabilities complicate capability assessment and safety testing; models may have abilities not anticipated at launch.

See: [AI safety](#); [Evaluation \(evals\)](#); [Scaling law](#)

## [Encoder](#)

T

A neural network component that compresses inputs into lower-dimensional representations. These compressed representations are used for retrieval and other tasks.

See: [Decoder](#); [Embedding](#); [Transformer](#)

## [Encoder-only model](#)

T

A transformer model using only the encoder, suited for classification and extraction tasks. Encoder-only models are used for classification, not generation.

See: [Classification](#); [Decoder-only model](#); [Transformer](#)

## [Encryption](#)

C

Cryptographic protection of data at rest and in transit. Encryption is baseline security for AI systems; clarify whether prompts, outputs, and logs are encrypted.

See: [Confidential computing](#); [Data security](#); [Security](#)

### [End-of-life / deprecation \(EOL\)](#)

The point at which a vendor or developer stops supporting a model, API, or product version (e.g., no updates, limited availability, retirement). Deprecation affects continuity, change control, and portability planning. Despite broad improvement in model capabilities, the deprecation of a model can change risks in a particular AI system for the worse.

See: [Business continuity](#); [Change control](#); [Model drift](#); [Portability](#)

### [Endpoint](#)

C

A network-accessible API path where requests are sent. Endpoint scope (public vs. private, authentication) is a key security factor.

See: [Authentication](#); [Security](#)

### [Environmental impact](#)

R

The energy consumption, carbon emissions, and resource use associated with AI training and inference. Environmental concerns increasingly appear in ESG reporting and procurement criteria; some regulations require disclosure.

See: [Compute](#); [Sustainability](#); [Training](#)

### [Epoch](#)

One complete pass through the entire training dataset during model training. Affects training cost and model behavior.

See: [Batch size](#); [Hyperparameter](#); [Training](#)

### [Ethical AI](#)

A broad term describing efforts to develop and use AI in ways aligned with stated values (e.g., fairness, transparency, accountability, privacy, safety). The term is used in governance frameworks and policy statements and is not a single technical standard.

See: [AI governance](#); [Fairness](#); [Responsible AI](#); [Transparency](#)

### [EU AI Act](#)

R

The European Union's primary AI regulation using risk-based categorization and obligations for providers and deployers. Key terms (provider, deployer, GPAI, high-risk, prohibited practices) affect compliance strategy for EU-facing deployments.

See: [Deployer](#); [GPAI](#); [High-risk AI system](#); [Prohibited AI practices](#); [Provider](#)

## Evaluation (evals)

C T

Repeatable testing measuring model/system quality, safety, robustness, and compliance with requirements. Evals support governance, regulatory compliance, and reasonable care; define who runs them and how results are handled.

See: [Benchmark](#); [Monitoring](#); [Red teaming](#); [Safety evaluation](#)

## Evasion attack

An adversarial attack crafting inputs to cause misclassification of malicious content as benign. Evasion attacks can bypass content moderation and security filters; relevant to security representations.

See: [Adversarial attack](#); [Content filtering](#); [Robustness](#)

## Excessive agency

R

A risk where AI systems are given more autonomy, permissions, or capabilities than necessary for their intended function. OWASP Top 10 for LLMs identifies excessive agency as a key vulnerability; least privilege principles apply.

See: [Agentic AI](#); [Least privilege](#); [Tool permissions](#)

## Execution

In the CORE framework, the dataflow connecting components from input to output, documenting how data travels through the system. Multiple execution paths may exist based on routing logic, conditional branches, or error handling. Documenting execution paths is needed for compliance with safety-by-design and privacy-by-design regulations as well as explainable AI.

See: [Audit](#); [Components](#); [CORE](#); [Logging](#); [Operations](#); [Resources](#); [XAI](#)

## Executive Orders on AI

R

A series of U.S. presidential executive orders addressing artificial intelligence policy, with significant shifts between administrations. Key orders include: President Trump's 2019 order on "Maintaining American Leadership in Artificial Intelligence" (later codified in the National AI Initiative Act of 2020); President Biden's October 2023

Executive Order 14110 on "Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence," which established reporting requirements for frontier models, mandated red-teaming, defined "dual-use foundation model," and directed creation of the AI Safety Institute; and President Trump's January 2025 orders revoking EO 14110 and replacing it with "Removing Barriers to American Leadership in Artificial Intelligence," which shifted policy toward deregulation and innovation. A December 2025 order sought to establish federal preemption of state AI laws. Despite the rescission of EO 14110, terminology it introduced (such as "dual-use foundation model" and "red team") remains in common use, and voluntary frameworks like the NIST AI RMF developed pursuant to it continue to be referenced in procurement and governance. Executive orders bind federal agencies but do not directly regulate private parties; however, they influence federal procurement requirements, agency enforcement priorities, and industry standards.

See: [Dual-use foundation model](#); [NIST AI RMF \(AI Risk Management Framework\)](#); [Red teaming](#)

### Exfiltration

Unauthorized extraction of sensitive data from a system, including via model outputs or tool calls. Exfiltration risk is central to confidentiality and privacy; address via least privilege, prompt injection defenses, and logging.

See: [Data leakage](#); [Prompt injection](#); [Tool permissions](#)

### Explainability

C R

The ability to describe why a model produced a particular output in understandable terms. Explainability is often requested in regulated decisions. Due to limited understanding of exactly which factors lead to a particular output, explainability may be low. "Thinking" traces from reasoning models may be helpful for understanding why a model could generate a particular output, but they do not accurately represent what actually went into any particular decision.

See: [Explainable AI](#); [Interpretability](#); [Right to explanation](#)

### Explainable AI (XAI)

Techniques and methods designed to make AI model behavior more understandable to humans. XAI encompasses various approaches with different fidelity-complexity tradeoffs; claims often specify which methods are used and their limitations. Be careful not to confuse with X.ai, the AI provider associated with the microblogging platform X.

See: [Explainability](#); [Interpretability](#)

### Export controls

R



Laws restricting export of certain technologies (including AI hardware and software) to certain countries or parties. Export controls can affect model/compute sourcing, cross-border hosting, and M\&A diligence.

See: [Compute](#); [Dual-use](#)

### [Extraction attack](#)

An adversarial attack seeking to extract training data, model weights, or system prompts through queries. Extraction attacks can expose confidential information and IP; relevant to security testing and rate limiting design.

See: [Data privacy attack](#); [Model extraction](#); [Prompt leakage](#)

## **F**

### [F1 score](#)

A metric combining precision and recall, useful when classes are imbalanced. If a vendor promises "performance," clarify which metric matters (accuracy vs. F1 vs. recall) and on what dataset.

See: [Accuracy](#); [Benchmark](#); [Precision](#); [Recall](#)

### [Fairness](#)

L R

The principle that AI systems should not produce unjustified differential outcomes for different groups. Fairness has multiple technical definitions (e.g. demographic parity, equalized odds, individual fairness) that can conflict with each other. A system cannot simultaneously satisfy all fairness definitions in most real-world scenarios. Contracts and governance documents should specify which definition applies and how it will be measured.

See: [Algorithmic discrimination](#); [Bias](#); [Disparate impact](#)

### [Fair use](#)

L

A U.S. copyright doctrine permitting certain uses of copyrighted material without permission. Fair use is central to AI training data litigation; the doctrine's application to AI is actively being litigated.

See: [Copyright infringement \(AI context\)](#); [Training data](#); [Transformative use](#)

### [False negative](#)

An instance incorrectly classified as negative when the true label is positive. False negative rates matter for safety-critical applications (missed fraud, missed medical conditions).

See: [False positive](#); [Precision](#); [Recall](#)

### [False positive](#)

An instance incorrectly classified as positive when the true label is negative. False positive rates affect user experience and can create liability (wrongful denials, false accusations).

See: [False negative](#); [Precision](#); [Recall](#)

### [Feature](#)

An individual measurable property of data used as input to a model (e.g., age, location, purchase history, medical codes). Feature choice affects model behavior; in regulated contexts, using protected characteristics (or close proxies) as features may trigger heightened review or restrictions depending on the use case and jurisdiction. In older models, features were frequently chosen by data scientists; in modern large models, features are usually discovered as part of the training process and may not be explicit.

See: [Bias](#); [Feature engineering](#); [Training data](#)

### [Feature engineering](#)

The process of selecting, transforming, and creating features for model training. Feature engineering choices can introduce bias or encode protected characteristics indirectly.

See: [Bias](#); [Feature](#); [Training](#)

### [Federated learning](#)

A training approach where model updates happen on distributed devices without centralizing raw data. Federated learning may be used as a privacy-enhancing technique.

See: [Data minimization](#); [Privacy-enhancing technology](#); [Training](#)

### [Feedback loop](#)

C

A cycle where model outputs influence future training data or model behavior. Feedback loops can convert inference data into training data; contracts often address whether and how feedback is used.

See: [Monitoring](#); [Service improvement](#)

## Few-shot prompting

T

Providing a small number of examples in the prompt to guide model behavior without changing weights. Few-shot prompts act like operational policy; treat prompt libraries as sensitive assets with change control.

See: [In-context learning](#); [Prompting](#); [Zero-shot prompting](#)

## Field of use restriction

C

A license term limiting where, how, or for what purposes a model may be used. Field-of-use restrictions are common in model licenses and can create compliance and enforcement issues.

See: [AUP](#); [License grant](#); [Open weights](#)

## Fine-tuned model

C

A model produced by fine-tuning a base model on specific data. Determine ownership, permitted reuse, and confidentiality rules for fine-tuned artifacts and training datasets.

See: [Adapter](#); [Fine-tuning](#); [Model artifact](#)

## Fine-tuning

C T

Additional training of a pre-trained model on specialized data to customize behavior. Fine-tuning raises the sharpest IP and confidentiality issues: data rights, artifact ownership, and training set retention.

See: [Adapter](#); [LoRA \(Low-Rank Adaptation\)](#); [Pre-training](#); [Transfer learning](#)

## Floating-point operations (FLOPs)

A measure of computational work; used to quantify training and inference costs. FLOP thresholds appear in regulatory definitions (e.g. the EU AI Act and AI Executive Orders); relevant to compute cost negotiations.

See: [Compute](#); [Training](#)

## Foundation model

C R T

A large, general-purpose model trained on broad data that can be adapted to many tasks. Foundation model sourcing and licensing affect compliance posture; clarify whether vendor provides proprietary, open-weights, or wrapper.

See: [Base model](#); [GPAI](#); [LLM](#); [Open weights](#)

## [Frontier model](#)

R

An AI model at or near the cutting edge of capabilities, typically characterized by training compute, parameter count, and emergent abilities that may pose novel risks not present in less capable systems. Regulatory frameworks have attempted to formalize this concept using compute thresholds: President Biden's EO 14110 defined "dual-use foundation model" at  $10^{26}$  FLOPs; the EU AI Act presumes "systemic risk" for GPAI models trained above  $10^{25}$  FLOPs. Though EO 14110 was rescinded, compute thresholds remain relevant to export controls, international coordination, and voluntary industry commitments. The threshold for "frontier" shifts continuously—today's frontier becomes tomorrow's baseline.

See: [Dual-use foundation model](#); [Executive Orders on AI](#); [Export controls](#); [Foundation model](#); [GPAI](#); [Systemic risk](#)

## [Function calling](#)

T

See: [Tool calling \(function calling\)](#)

## [Fundamental Rights Impact Assessment \(FRIA\)](#)

R

An assessment required under the EU AI Act for deployers of high-risk AI systems used to evaluate individuals, examining potential impacts on fundamental rights including privacy, non-discrimination, human dignity, and access to justice. FRIAs must be completed before first use and updated when circumstances materially change. Unlike DPIAs (which focus on data protection), FRIAs address broader rights impacts and require consideration of specific affected populations.

See: [Deployer](#); [EU AI Act](#); [High-risk AI system](#); [Risk assessment](#)

# **G**

## [General Data Protection Regulation \(GDPR\)](#)

R

The EU's comprehensive data protection framework governing processing of personal data. In AI, GDPR concepts commonly implicated include lawful basis, transparency, data minimization, purpose limitation, security, cross-border transfers, and rights (e.g., access, deletion, objection); applicability depends on the processing context and the roles of controller/processor.

See: [Automated decision-making](#); [Controller / processor](#); [Personal data](#)

## Generalization

A model's ability to perform well on new, unseen data beyond its training set. Generalization relates to reliability; poor generalization can undermine product claims and cause failure on edge cases.

See: [Evaluation \(evals\)](#); [Overfitting](#); [Robustness](#)

## General-purpose AI (GPAI)

R

AI designed for broad applicability across many tasks and contexts; a category in the EU AI Act. GPAI terminology appears in EU AI Act with specific obligations for GPAI providers, especially those with systemic risk.

See: [EU AI Act](#); [Foundation model](#); [Systemic risk](#)

## Generated content

Content produced by a generative model (text, images, audio, video, code). “Generated content” is often used interchangeably with “Output,” though contracts may define these differently.

See: [Generative AI \(GenAI\)](#); [Output](#)

## Generation

C

Using a trained model to process new inputs and create new, unknown outputs. Unlike classification, recommendation, or extraction systems, generation is not constrained to predefined categories; the space of possible outputs is effectively unbounded. Generative responses will frequently include information provided as part of the prompt or from the training inputs.

See: [Generative AI \(GenAI\)](#); [Hallucination](#); [Inference](#)

## Generative Adversarial Network (GAN)

T

A generative model architecture using competing generator and discriminator networks. GANs powered early deepfake technology; largely superseded by diffusion models but still relevant for understanding synthetic media.

See: [Deepfake](#); [Generative AI \(GenAI\)](#); [Synthetic data](#)

## Generative AI (GenAI)

C L

Models that generate new content (text, code, images, audio, video) rather than only classifying inputs. Generative outputs raise distinctive risks: hallucination, defamation, IP infringement, confidentiality leakage, deepfakes.

See: [Diffusion model](#); [Hallucination](#); [LLM](#)

## [Governance, Risk, and Compliance \(GRC\)](#)

C

Term describing frameworks for managing governance policies, enterprise risk management, and regulatory compliance. Traditional GRC relies on organizational controls (such as written policies, training, procedures, and attestations) that work because humans read, understand, and follow them. In contrast, written policies do not constrain an AI model; the policy must be translated into technical controls such as guardrails, system prompts, tool permissions, and monitoring that govern actual system behavior. Effective AI governance requires mapping organizational controls (which govern humans who build and oversee AI) to technical controls (which govern what AI systems can do). Organizations with mature GRC functions can accelerate AI governance, but AI governance requires enforcement mechanisms beyond those designed for human compliance.

See: [AI governance](#); [Controls](#); [CORE](#); [Guardrails](#); [Risk assessment](#)

## [GPAI](#)

See: [General-purpose AI](#)

## [Gradient descent](#)

An optimization method that adjusts model weights in small steps toward lower prediction error. Think of it as rolling a ball downhill: the algorithm repeatedly moves weights in whatever direction reduces the loss function. Gradient descent determines both training speed and whether the model converges on useful patterns.

See: [Backpropagation](#); [Loss function](#); [Training](#)

## [Graphics Processing Unit \(GPU\)](#)

C T

Specialized hardware used to accelerate AI training and inference. GPU availability affects costs, vendor lock-in, and export control considerations.

See: [Compute](#); [Export controls](#)

## [Groundedness](#)

A RAG quality metric measuring whether outputs are supported by the provided context rather than fabricated. Groundedness is key to RAG reliability claims; ungrounded outputs may involve fabrication.

See: [Grounding](#); [Hallucination](#)

## [Grounding](#)

C T

Constraining outputs to specified sources (retrieved documents, databases) rather than model patterns alone. Grounding improves defensibility and reduces hallucination risk; requires controls over source corpus and retrieval logs.

See: [Citation](#); [Hallucination](#)

## [Ground truth](#)

The best available correct labels or answers used to train or evaluate a system. Ground truth can be disputed in subjective domains; define labeling standards and dispute processes.

See: [Annotation](#); [Evaluation \(evals\)](#); [Labeling](#)

## [Group Relative Policy Optimization \(GRPO\)](#)

A training technique for reasoning models that optimizes based on relative performance within groups of responses. Introduced in late 2024 and used in many frontier models.

See: [Direct Preference Optimization \(training\)](#); [Reasoning model](#)

## [Guardrails](#)

C

Technical and procedural controls limiting unsafe or undesired behavior. Guardrails are often the practical basis for safety representations and compliance strategies.

See: [Content filtering](#); [HITL](#); [Safety policy](#)

# **H**

## [Hallucination](#)

C L T

Plausible-sounding output that is wrong, unsupported, or fabricated, including fake citations. Hallucinations drive malpractice, consumer protection, and litigation risk; mitigation combines grounding, evals, and human review.

See: [Grounding](#); [Verification](#)

## [High-risk AI system](#)

R

An AI system category with heightened regulatory obligations due to impact on safety or fundamental rights. High-risk categorization drives documentation, testing, human oversight, and post-market monitoring duties under the EU AI Act.

See: [Conformity assessment](#); [EU AI Act](#); [Monitoring](#)

## [HITL](#)

See: [Human-in-the-loop](#)

## [Hosting](#)

### C

The environment where a model or AI system runs (cloud, on-prem, managed service), including compute, networking, and storage. Hosting affects security controls, data residency, latency, and incident response responsibilities.

See: [Cloud provider](#); [Data residency](#); [On-prem deployment](#); [Security](#)

## [Human evaluation](#)

### T

Assessment of AI outputs by human reviewers for quality, safety, or compliance. Document human evaluation methods to support quality claims and manage reviewer exposure to sensitive content.

See: [Evaluation \(evals\)](#); [Labeling](#); [Red teaming](#)

## [Human-in-the-loop \(HITL\)](#)

### C R

A control where a human reviews, approves, or edits outputs before use. HITL can be a decisive risk mitigator and regulatory requirement for high-stakes decisions.

See: [Approval workflow](#); [Automated decision-making](#); [Autonomy level](#)

## [Hybrid search](#)

Retrieval that combines traditional keyword search with vector/semantic search to improve recall and precision. Hybrid search is common in enterprise RAG implementations.

See: [Semantic search](#); [Vector database](#)

## [Hyperparameter](#)

### T

A configuration setting chosen by engineers (not learned), such as learning rate or temperature. Hyperparameters affect performance and reproducibility; relevant in disputes about model version changes.



See: [Parameter](#); [Temperature](#); [Training](#)

## I

### [Image generation](#)

AI creation of images from text prompts or other inputs. Image generation raises copyright questions regarding training data and outputs, deepfake and NCII concerns, and trademark issues when generated images depict recognizable brands or logos.

See: [Diffusion model](#); [Generative AI \(GenAI\)](#); [Text-to-image](#)

### [Incident response](#)

## C

Processes for detecting, responding to, and recovering from AI system failures or security events. AI incident response plans often address model-specific scenarios such as hallucination-caused harm, prompt injection breaches, and unexpected behavior changes; contracts commonly specify notification and cooperation obligations.

See: [Business continuity](#); [Monitoring](#); [Security](#)

### [In-context learning](#)

A model's ability to adapt behavior within a prompt using provided examples, without weight changes. Because in-context learning allows prompts and retrieved documents to materially change behavior, prompt content is often treated as part of the controlled system.

See: [Context window](#); [Few-shot prompting](#); [Prompting](#)

### [Indemnity](#)

## C

A contractual promise by one party to defend and/or reimburse the other for specified third-party claims (often including costs). In AI contexts, indemnities commonly address IP claims, data protection incidents, and misuse claims, with scope turning on definitions of “Input,” “Output,” and “Training” and on compliance with use restrictions.

See: [IP indemnity](#); [Limitation of liability](#); [Warranty](#)

### [Indirect prompt injection](#)

## T

A prompt-injection technique where malicious instructions are embedded in data the system retrieves or processes (e.g., web pages, emails, tickets, documents) rather than

in the user's direct prompt. This is a common risk in RAG and tool-enabled agent systems.

See: [Connector](#); [Prompt injection](#); [Tool calling \(function calling\)](#)

## [Inference](#)

C T

Using a trained model to process new inputs and create outputs. This is the operational phase, as distinguished from training. Inference is distinguished from training by the way in which data is handled: training uses input data to adjust model weights, leading to possible issues with memorization, while inference only uses input data transiently to produce outputs. Inference is also distinguished from generation; a classification system performs inference, but the scope of possible outputs is limited. In contrast, generative systems may reproduce input or training data. Contracts often define "inference" and "training" separately and impose different restrictions on each.

See: [Generation](#); [Latency](#); [Training](#)

## [Inference-time scaling](#)

T

Computational resources used during inference, particularly the extended processing in reasoning models that "think" before responding. Inference-time compute affects cost and latency; reasoning models may use substantially more compute per query than standard models.

See: [Compute](#); [Inference](#); [Reasoning model](#)

## [Input Data](#)

C

Content provided to an AI system as input, including prompts and context documents. Input Data often contains sensitive information, so contracts commonly define whether it can be retained, logged, or used for service improvement.

See: [Confidential information](#); [Logging](#); [User prompt](#)

## [Insecure output handling](#)

A security risk where AI outputs are used in downstream systems without adequate validation or escaping (e.g., injecting generated content into code, HTML, SQL, or commands). This can convert hallucinations or malicious outputs into system actions.

See: [Prompt injection](#); [Security](#); [Tool calling \(function calling\)](#)

## [Instruction hierarchy \(prompt precedence\)](#)

T

The priority order that determines which instructions a model follows when multiple sources exist (commonly: system prompt > developer instructions > user prompt > tool outputs/retrieved text). Understanding precedence is central to prompt injection and guardrail design.

See: [Developer prompt](#); [Prompt injection](#); [System prompt](#); [User prompt](#)

### [Instruction tuning](#)

Fine-tuning a model to follow instructions and engage in dialogue. Instruction-tuned models (often called "chat" models) behave differently than base models, which is relevant when assessing capabilities and limitations.

See: [Chat model](#); [Fine-tuning](#)

### [Intellectual property \(IP\)](#)

C L

Legal rights in creations of the mind, including patents, copyrights, trademarks, and trade secrets. AI raises novel IP questions across all categories, with training data, outputs, and model architecture each presenting distinct issues requiring separate analysis.

See: [Copyright](#); [IP indemnity](#); [Patent](#); [Trade secret](#)

### [Intended purpose](#)

R

Under the EU AI Act, the use for which an AI system is intended by the provider as specified in required documentation, including instructions for use, promotional materials, and technical specifications.

See: [EU AI Act](#); [High-risk AI system](#); [Off-label use](#); [Provider](#); [Technical documentation](#)

### [Intended use](#)

C R

The purpose, context, and conditions under which a model or AI system is designed and evaluated to operate (e.g., internal drafting vs automated decisions; healthcare vs general productivity). "Intended use" is commonly used to scope warranties, safety controls, and regulatory obligations.

See: [AUP](#); [Evaluation \(evals\)](#); [Off-label use](#); [Risk assessment](#)

### [Interpretability](#)

The degree to which a human can understand how a model produces its outputs. Interpretability is stronger than explainability and implies genuine understanding of

internal mechanisms; truly interpretable models are often less capable than black-box alternatives.

See: [Black box](#); [Explainability](#); [XAI](#)

## IP

See: [Intellectual property](#)

## IP indemnity

### C

A contractual promise to defend and compensate for intellectual property infringement claims. AI IP indemnities vary widely in scope (training data vs. output), exceptions (user modifications, combinations), and conditions (cooperation, control of defense), requiring careful negotiation.

See: [Copyright](#); [Indemnity](#); [Training data](#)

## ISO/IEC 27001

### C

An international standard for information security management systems (ISMS). References to ISO 27001 often appear in security questionnaires and vendor contracts as evidence of a structured security program.

See: [Security](#); [Security addendum](#); [SOC 2](#)

## ISO/IEC 42001

### C R

An international standard for AI management systems (AIMS), ISO 42001 follows the Plan-Do-Check-Act structure of other ISO management system standards (like ISO 27001 for information security), making it integrable with existing compliance programs. The standard covers AI system lifecycle management, risk assessment and impact evaluation, data governance, third-party supplier oversight, and 38 specific controls in its annexes. Organizations already certified to ISO 27001 can leverage significant structural overlap.

See: [AI governance](#); [Audit](#); [ISO/IEC 27001](#); [NIST AI RMF \(AI Risk Management Framework\)](#); [Risk assessment](#); [SOC 2](#)

## J

## Jailbreak

### C

An attempt to bypass a model's safety restrictions through crafted prompts. Jailbreak resistance is part of safety evaluation, and successful jailbreaks may trigger AUP violations, incident response obligations, and potential liability for resulting harms.

See: [Adversarial attack](#); [Prompt injection](#); [Safety evaluation](#)

## Jamba

A family/name used for architectures combining elements of transformers with state space models (SSMs) to improve efficiency for long sequences. The term is used in technical discussions about model architecture and inference cost.

See: [Context window](#); [Transformer](#)

## JSON mode

A setting constraining model outputs to valid JSON format. Structured output modes like JSON mode improve reliability and parsing for automated workflows, reducing integration errors and enabling programmatic processing of AI outputs.

See: [Function calling](#); [Structured output](#)

# K

## Key-Value cache (KV)

Memory storing prior context during inference to enable efficient generation. KV cache size affects context window limits and inference cost, which is relevant when understanding capacity constraints and pricing models.

See: [Context window](#); [Memory](#); [Transformer](#)

## Knowledge base

# C

A curated collection of documents or data used for retrieval in RAG systems. Knowledge base contents determine output quality and risk exposure, making access control, accuracy verification, and update processes important governance considerations.

See: [Connector](#); [Grounding](#)

## Knowledge distillation

# T

See: [Distillation](#)

## Knowledge transfer

The process of moving know-how, documentation, and operational understanding from one team or vendor to another (e.g., during vendor transitions, M\&A, or outsourcing). In AI deployments, knowledge transfer may include model documentation, evals, runbooks, and data pipelines.

See: [Business continuity](#); [Documentation](#); [Portability](#)

## L

### Labeling

See: [Annotation](#)

### Language model

#### T

A model that assigns probabilities to sequences of tokens; modern large language models (LLMs) are language models scaled and trained for broad capabilities. "Language model" can also refer to smaller or domain-specific models.

See: [Large Language Model](#); [Token](#); [Transformer](#)

### Large Language Model (LLM)

#### C T

A neural network trained on extensive text data to understand and generate language. "LLM" is often used interchangeably with "AI" in business contexts, though technically it refers to text-focused models; modern LLMs are increasingly multimodal.

See: [Foundation model](#); [Generative AI \(GenAI\)](#); [Small Language Model](#); [Transformer](#)

### Large Reasoning Model (LRM)

An LLM specifically trained for complex reasoning, often using extended chain-of-thought processing. Reasoning models like o1 and DeepSeek-R1 represent a capability shift and may be more reliable for complex analysis, though they have different cost and latency profiles.

See: [Chain of Thought \(CoT\)](#); [Inference-time scaling](#); [Reasoning model](#)

### Latency

#### C T

The time delay between sending a request and receiving a response. Latency affects user experience and may be specified in SLAs; high latency can undermine real-time use cases and is often tested under realistic load conditions.

See: [Inference](#); [Service Level Agreement / Service Level Objective](#); [Throughput](#)

## Latent space

The abstract multi-dimensional space encoding learned representations of data. Embeddings exist in latent space, and understanding this concept helps explain how AI systems represent and compare semantic meaning.

See: [Embedding](#); [Representation learning](#)

## Lawful basis

R

Under GDPR, one of six legal grounds that must exist before processing personal data: consent, contractual necessity, legal obligation, vital interests, public task, or legitimate interests. AI systems often rely on legitimate interests (requiring a balancing test) or consent (requiring clear, specific, freely-given agreement). Training on personal data, inference processing, and service improvement uses may each require separate lawful basis analysis.

See: [Consent](#); [Personal data](#); [Privacy](#); [Purpose limitation](#)

## Least privilege

C

The security principle of granting only the minimum permissions necessary for an actor (user, service, or agent) to perform its task. In agentic and tool-enabled systems, least-privilege permissioning and scoped tool access are common controls to reduce the impact of errors, abuse, or prompt injection.

See: [Access control](#); [Security](#); [Tool permissions](#)

## Lethal Trifecta

A security vulnerability pattern identified by Simon Willison occurring when an AI agent simultaneously possesses three capabilities: (1) access to private or sensitive data, (2) exposure to untrusted content, and (3) the ability to communicate externally. When all three capabilities are present, prompt injection attacks can cause the agent to access private data and transmit it to an attacker. The Lethal Trifecta has been demonstrated against major products including Microsoft 365 Copilot, ChatGPT, Google Gemini, Slack, and GitHub Copilot. Because prompt injection remains an unsolved problem, the primary defense is to ensure AI systems never combine all three capabilities simultaneously.

See: [Agentic AI](#); [Agents Rule of Two](#); [Exfiltration](#); [Prompt injection](#); [Tool permissions](#)

## License compatibility

Whether multiple licenses (e.g., open source licenses, model licenses, and proprietary licenses) can be complied with simultaneously when components are combined or

distributed. Incompatibilities can arise from obligations such as copyleft, attribution, field-of-use limits, or downstream restrictions.

See: [Copyleft license](#); [Open source](#)

## [License compliance](#)

C

Adhering to terms of software and model licenses. AI systems often combine multiple licensed components with different terms, requiring tracking and satisfying all applicable obligations to avoid infringement claims.

See: [Copyleft license](#); [Open source](#)

## [License grant](#)

C

The specific permissions conveyed by a license. Model license grants vary widely in scope, field of use restrictions, sublicensing rights, and modification permissions, requiring careful examination for each intended use.

See: [Field of use restriction](#); [IP](#); [Open weights](#)

## [Limitation of liability](#)

C

Contract terms that limit damages (e.g., caps, exclusions of consequential damages, and carve-outs). In AI service agreements, parties often allocate risk differently across categories such as IP claims, security incidents, confidentiality, and misuse; the negotiated structure varies by use case and regulatory exposure.

See: [Contract](#); [Indemnity](#); [Warranty](#)

## [Litigation hold](#)

L

A process to preserve relevant information when litigation or an investigation is reasonably anticipated. In AI systems, holds may apply to logs, prompts, outputs, tool call records, and model/version artifacts.

See: [E-discovery](#); [Logging](#); [Retention](#)

## [LLM](#)

T

See: [Large Language Model](#)

## [Logging](#)

C



Recording system activity including prompts, outputs, and operational data. Logging is essential for audit and debugging but creates privacy and confidentiality exposure; contracts commonly define what is logged, who can access logs, and retention periods.

See: [Monitoring](#); [Records retention](#); [Usage data / telemetry](#)

## [LoRA \(Low-Rank Adaptation\)](#)

C T

A parameter-efficient fine-tuning technique that trains small adapter weights rather than full model weights. LoRA adapters are common deliverables in custom AI projects, so ownership, portability to other base models, and confidentiality terms are often stated explicitly.

See: [Adapter](#); [Fine-tuning](#); [Parameter-efficient fine-tuning](#)

## [Loss function](#)

A mathematical formula that measures how wrong the model's predictions are. The choice of loss function determines what the model optimizes for; a model trained to minimize one type of error may perform poorly by other measures. This is relevant when evaluating whether a model was designed appropriately for its intended use: a model optimizing for average accuracy may systematically fail on minority cases.

See: [Gradient descent](#); [Training](#)

## [LRM](#)

See: [Large Reasoning Model](#)

# **M**

## [Machine learning \(ML\)](#)

A subset of AI where systems learn patterns from data rather than following explicit rules. ML is the technical foundation of modern AI, and understanding that models learn statistical patterns (rather than "knowing" facts) helps assess capabilities and limitations.

See: [Deep learning](#); [Neural network](#); [Training](#)

## [Machine Unlearning](#)

R

Techniques intended to reduce or remove the influence of specific data from a trained model without full retraining. Approaches and effectiveness vary; some providers offer deletion workflows or re-training processes, but the scope of guarantees, documentation, and verification methods differs across implementations.

See: [Algorithmic Disgorgement \(Model Deletion\)](#); [Guardrails](#); [Right to be Forgotten](#)

## [Mamba](#)

T

A selective state space model architecture offering efficient processing of long sequences. Mamba represents an alternative to transformer architecture for some applications, particularly those requiring very long context windows.

See: [Jamba](#); [Transformer](#)

## [Massive Multitask Language Understanding \(MMLU\)](#)

A benchmark testing model knowledge across many academic subjects. MMLU scores are commonly cited in model marketing, but benchmarks don't guarantee real-world performance and may not reflect capabilities on domain-specific tasks.

See: [Benchmark](#); [Evaluation \(evals\)](#)

## [MCP](#)

See: [Model Context Protocol](#)

## [Membership inference](#)

T

An attack determining whether specific data was used in training. Membership inference can reveal training data composition, which is relevant to privacy claims, trade secret disputes, and copyright litigation discovery.

See: [Data privacy attack](#); [Memorization](#); [Model inversion](#)

## [Memorization](#)

C L

A model's tendency to reproduce training data verbatim rather than generalizing. Memorization creates copyright and privacy exposure and is the mechanism behind extraction attacks; the degree of memorization varies by model and data frequency. Usually the result of overfitting to repeatedly-seen training data.

See: [Copyright](#); [Data leakage](#); [Extraction attack](#); [Overfitting](#)

## [Memory](#)

Stored state used across interactions (e.g., conversation history, user preferences, task state). Memory can be ephemeral (within a context window) or persistent (stored and retrieved later), and it can raise retention, privacy, and confidentiality considerations.

See: [Context window](#); [Logging](#); [Personal data](#); [Retention](#)

## Metadata

Data describing other data, including creation dates, sources, and processing history. AI-generated content may lack authentic metadata or have synthetic metadata, which is relevant to evidence authentication and content provenance disputes.

See: [Content provenance](#); [Data provenance](#)

## Mixture of Experts (MoE)

T

A neural network architecture where only a subset of the model's parameters, called "experts", are activated for each input. MoE allows models to have very large total parameter counts while keeping inference costs manageable: a model with 400 billion parameters might activate only 50 billion for any given query.

See: [Architecture](#); [Compute](#); [Inference](#); [Parameter](#)

## MLOps

Practices for deploying and maintaining ML systems in production, including monitoring, versioning, and updates. MLOps maturity affects reliability, reproducibility, and change control; assess vendor MLOps practices as part of due diligence.

See: [DevOps](#); [Model registry](#); [Monitoring](#)

## Model

C

The trained artifact (weights and architecture) that processes inputs to produce outputs. In contracting, "model" is commonly distinguished from "AI system" because model licensing and ownership are often negotiated separately from system-level concerns like hosting, data handling, and integration.

See: [AI system](#); [Model artifact](#); [Weights](#)

## Model artifact

C

Tangible outputs of model development and deployment, such as weights, checkpoints, adapters, fine-tuned models, training logs, evaluation results, and prompt templates. These artifacts are often treated as valuable intellectual property and can be addressed in development, licensing, and confidentiality terms.

See: [Adapter](#); [Checkpoint](#); [Weights](#)

## Model capability

The range and level of tasks a model can perform reliably, often evidenced through evaluations, benchmarks, and red teaming results. Capability statements can affect intended-use scoping and governance classification.

See: [Benchmark](#); [Capabilities](#); [Evaluation \(evals\)](#); [Intended use](#)

### Model card

Structured documentation describing a model's intended use, performance characteristics, and limitations. Model cards support due diligence and can undermine marketing claims if disclosed limitations conflict with vendor representations.

See: [Documentation](#); [Evaluation \(evals\)](#); [Technical documentation](#)

### Model collapse

C

The degenerative process where a model trained on synthetic (AI-generated) data eventually loses quality, diversity, and connection to reality. As the internet fills with AI-generated content, "organic" human-generated training data becomes a premium asset; contracts may need to specify the ratio of synthetic versus organic data to support data quality warranties.

See: [Data provenance](#); [Synthetic data](#); [Training data](#)

### Model compression

Techniques reducing model size while maintaining performance, including quantization and distillation. Compressed models may behave differently than originals, and compression is a form of modification that may require license analysis.

See: [Distillation](#); [Edge deployment](#); [Quantization](#)

### Model Context Protocol (MCP)

An open source standard protocol for connecting AI models to external data sources and tools. MCP addresses the execution layer: how an agent calls tools and retrieves data. MCP enables interoperability between different AI systems and data sources, with implications for data access control, logging, and vendor lock-in. The use of MCP (usually described in terms of an “MCP server”) almost always implies that an AI system will be given access to some Resource.

See: [Connector](#); [CORE](#); [Resources](#); [Tool calling \(function calling\)](#)

### Model deprecation

See: [End-of-life / deprecation](#)

### Model drift

Changes in model behavior over time due to updates, data changes, or environmental shifts. Drift drives change control clauses in contracts and can create unexpected compliance failures in validated workflows.

See: [Change control](#); [Monitoring](#); [Version pinning](#)

### [Model escrow / artifact escrow](#)

C

An arrangement where model weights, training data, adapters, system prompts, or other AI artifacts are deposited with a third-party escrow agent for release to the customer upon specified trigger events.

See: [Adapter](#); [Business continuity](#); [Trade secret](#); [Vendor lock-in](#); [Weights](#)

### [Model extraction](#)

Attempting to recreate a model or approximate its behavior through repeated API queries. Model extraction is often prohibited in AUPs and relevant to trade secret protection; rate limiting and query monitoring serve as partial defenses.

See: [Distillation](#); [Extraction attack](#); [Rate limiting](#); [Trade secret](#)

### [Model improvements](#)

C

Enhancements to models developed during or after initial deployment, including new weights, adapters, and prompts. Ownership of improvements is a recurring dispute, particularly when customer data, feedback, or funding contributed to the work.

See: [Adapter](#); [Fine-tuned model](#); [License grant](#)

### [Model inversion](#)

Attacks inferring training data characteristics from model outputs. Model inversion is relevant to privacy assurances and security controls when sensitive data may have influenced training, even indirectly.

See: [Data privacy attack](#); [Membership inference](#); [Memorization](#)

### [Model license](#)

C

The terms governing how a model (weights, API access, or hosted service) may be used, including restrictions, attribution, field-of-use limits, and downstream obligations. Model licenses differ from open source software licenses and can be bespoke.

See: [Field of use restriction](#); [License grant](#); [Open source](#); [Open weights](#)

### [Model registry](#)

A system tracking model versions, metadata, approvals, and deployment status. Registries support governance and auditability by providing a single source of truth for what models exist, where they're deployed, and who approved them.

See: [Change control](#); [MLOps](#); [Model drift](#)

### [Model Risk Management \(MRM\)](#)

C R

A governance discipline for validating, monitoring, and controlling model use, particularly common in financial services under SR 11-7 guidance. MRM programs often impose documentation, independent validation, and change control obligations that apply to AI systems.

See: [AI governance](#); [Evaluation \(evals\)](#); [Monitoring](#)

### [Model supply chain](#)

The set of upstream components and processes used to build and operate a model or AI system (datasets, code, weights, third-party models, tools, connectors, hosting, and subprocessors). Supply chain analysis is used for security, IP provenance, and compliance.

See: [Security](#); [Subprocessor](#); [Training data](#)

### [Model update](#)

C

A change to model weights, tuning, safety settings, prompts, or configuration. Updates can introduce drift or unexpected behavior changes; contracts may require notice, version pinning options, and regression testing for critical workflows.

See: [Change control](#); [Model drift](#); [Version pinning](#)

### [Moderation](#)

C

Processes (automated and human) detecting and managing disallowed content or behavior. Moderation intersects with platform liability, AUP enforcement, and employment law; it also creates recordkeeping requirements and content reviewer welfare considerations.

See: [AUP](#); [Content filtering](#); [Safety policy](#)

### [Monitoring](#)

C R

Ongoing observation of a model or AI system in production (e.g., quality, safety events, latency, errors, drift, abuse signals). Monitoring outputs are used in incident response, governance reporting, and contract performance management.

See: [Data drift](#); [Incident response](#); [Model drift](#); [SLA/SLO](#)

### [Multi-agent system \(MAS\)](#)

An architecture where multiple AI agents collaborate or compete to accomplish tasks. Multi-agent systems create complex liability and attribution challenges because harm may result from emergent interactions rather than any single agent's action.

See: [Agentic AI](#); [Orchestration](#)

### [Multimodal](#)

T

See: [Multimodal model](#)

### [Multimodal model](#)

C T

A model accepting and/or producing multiple content types (text, images, audio, video). Multimodal capabilities expand privacy and IP risk through processing of faces, voices, and biometrics, and may trigger additional regulatory obligations.

See: [Biometric data](#); [Computer vision](#); [LLM](#)

### [Multi-tenancy](#)

C

Cloud architecture where multiple customers share infrastructure with logical separation. Multi-tenancy raises confidentiality concerns for AI systems; confirm isolation controls for prompts, embeddings, fine-tuned models, and logs.

See: [Data segregation](#); [SaaS \(Software as a Service\)](#); [Security](#)

### [Multi-tenant](#)

A service architecture in which multiple customers share underlying infrastructure and software while remaining logically separated. Multi-tenancy affects data segregation, logging, and security controls.

See: [Access control](#); [Data segregation](#); [SaaS \(Software as a Service\)](#)

## **N**

### [Narrow AI](#)

AI designed for specific tasks rather than general intelligence. All current AI systems are narrow AI, regardless of marketing claims; this is relevant when assessing vendor capability representations.

See: [Foundation model](#)

## [Natural Language Processing \(NLP\)](#)

AI techniques for understanding, analyzing, and generating human language. NLP encompasses many AI applications from classification to generation, with legal issues varying significantly by task type.

See: [Language model](#); [LLM](#); [Text classification](#)

## [NCII](#)

See: [Non-Consensual Intimate Images](#)

## [Neural network](#)

T

A computational architecture that processes information through layers of interconnected nodes, with each connection carrying a learned weight. Neural networks are not deterministic rule-based systems. Rather, they learn statistical patterns from training data, which explains both their capabilities and their tendency to produce confident-sounding errors. The term "neural" is a historical metaphor; these systems do not function like biological brains.

See: [Deep learning](#); [Transformer](#); [Weights](#)

## [NIST AI RMF \(AI Risk Management Framework\)](#)

R

A voluntary framework published by the National Institute of Standards and Technology providing guidance for managing AI risks throughout the system lifecycle. The framework is organized around four core functions: Govern (establishing accountability, policies, and culture), Map (understanding context, stakeholders, and potential impacts), Measure (assessing and tracking risks through evaluation and monitoring), and Manage (prioritizing and responding to identified risks). The framework is widely referenced in U.S. federal procurement, various AI Executive Orders, sector guidance, and enterprise customer requirements. NIST also published the Generative AI Profile (AI RMF 600-1) addressing risks specific to generative AI systems. Alignment with the NIST AI RMF is frequently requested in vendor assessments and can support reasonable care arguments, though "alignment" is self-assessed and does not involve certification or audit.

See: [AI governance](#); [Executive Orders on AI](#); [ISO/IEC 42001](#); [Risk assessment](#); [Trustworthy AI](#)



## NLP

See: [Natural Language Processing](#)

## Non-Consensual Intimate Images (NCII)

L R

Intimate images shared or generated without consent. AI-generated NCII is addressed by a growing set of criminal, civil, and platform policy regimes; obligations and remedies vary by jurisdiction and may depend on intent, distribution, and notice-and-takedown processes.

See: [Content filtering](#); [Deepfake](#); [Prohibited AI practices](#)

## Non-deterministic

A system property where the same input may produce different outputs across runs. Most generative AI systems are non-deterministic by default due to sampling strategies, floating-point computation variations, and infrastructure differences. Non-determinism affects auditability, testing reproducibility, and user expectations; controls like temperature settings can enforce more deterministic behavior when needed.

See: [Deterministic](#); [Reproducibility](#); [Sampling](#); [Temperature](#)

## Notice

C R

A communication provided to another party or to individuals (e.g., privacy notices, product disclosures, contractual notices of changes or incidents). “Notice” requirements are often defined by contract or applicable law, and may include timing and content requirements.

See: [Disclosure](#); [Incident response](#); [Transparency](#)

## No training on our data

C

A procurement requirement that a provider not use customer content (e.g., prompts, files, outputs, or connected data) to train, fine-tune, or otherwise improve its models beyond providing the contracted service. Implementations vary and may distinguish between model training, human review, debugging, safety monitoring, and logging/retention practices.

See: [Service improvement](#); [Training](#); [Usage data / telemetry](#)

# O

## Observability

## C

The ability to understand system behavior through logs, metrics, and traces. Observability supports audit, debugging, and incident response; contracts commonly define what telemetry is available and whether customers can access it.

See: [Incident response](#); [Logging](#); [Monitoring](#)

### Off-label use

Using an AI system for purposes beyond its intended or permitted use. Off-label use can affect contractual rights (including warranties and indemnities), compliance posture, and safety assumptions because the system may not have been evaluated or controlled for the new context.

See: [AUP](#); [Field of use restriction](#); [Intended use](#)

### On-prem deployment

## C

Running AI systems on customer-controlled infrastructure rather than vendor cloud. On-prem deployment affects data residency compliance, security responsibility allocation, update mechanisms, and support arrangements.

See: [Data residency](#); [Edge deployment](#); [SaaS \(Software as a Service\)](#)

### Open source

## C R

A term most precisely applied to software code distributed under an open source license. In AI discussions, “open source” is sometimes used more loosely to refer to publicly available weights, datasets, or systems with permissive access, which may not match established open source definitions.

See: [Model license](#); [Open source software](#); [Open weights](#)

### Open source software (OSS)

## C

Software distributed under licenses permitting use, modification, and redistribution with varying conditions. AI systems often incorporate open source components; compliance requires tracking all licenses and satisfying their respective obligations.

See: [Copyleft license](#); [License compliance](#)

### Open weights

## C

Models where trained weights are publicly released, but not necessarily under open source licenses. “Open weights” is distinct from “open source”; many open weight

licenses include field of use restrictions and may not meet the Open Source Initiative definition.

See: [Field of use restriction](#); [Foundation model](#); [License grant](#)

## Operations

In the CORE framework, the actions that components perform on external Resources or on data flowing through an AI system. Operations include data transformations (summarization, classification, generation), resource interactions (reading from databases, calling external APIs, writing outputs), and control functions (filtering, routing, logging).

See: [Components](#); [CORE](#); [Execution](#); [Resources](#); [Tool calling \(function calling\)](#)

## Optimization

Methods to improve a model or system's performance, cost, latency, or resource use (e.g., quantization, caching, batching, prompt compression). Optimization choices can affect accuracy, safety behavior, and reproducibility.

See: [Key-Value cache](#); [Latency](#); [Quantization](#)

## Orchestration

Coordinating multiple AI components, tools, and workflows to accomplish complex tasks. Orchestration layers make consequential decisions about tool selection and sequencing, requiring oversight, logging, and clear accountability.

See: [Agentic AI](#); [Multi-agent system](#); [Tool calling \(function calling\)](#)

## Output

### C

Content produced by an AI system in response to inputs. Outputs may contain confidential information from inputs or retrieval, create IP ownership questions, or cause harm through inaccuracy; contracts commonly define ownership, permitted uses, and retention.

See: [Generated content](#); [Hallucination](#); [IP](#)

## Overfitting

When a model performs well on training data but poorly on new data because it memorized specific examples rather than learning generalizable patterns. Overfitting explains why demo performance may not match production results, and frequently leads to memorization and subsequent regeneration of training material.

See: [Evaluation \(evals\)](#); [Generalization](#); [Memorization](#); [Training](#)

## OWASP Top 10 for LLMs

### R

A security framework identifying top vulnerabilities in LLM applications, including prompt injection, insecure output handling, and excessive agency. The OWASP Top 10 provides a standard reference for AI security requirements and risk assessment.

See: [Excessive agency](#); [Prompt injection](#); [Security](#)

## P

### Parameter

#### T

A learned value in a neural network that influences outputs; parameter count (often in billions) indicates model size. Parameter count is often cited as a capability proxy, but actual performance depends on architecture, training data, and post-training, not just size.

See: [Hyperparameter](#); [Model](#); [Weights](#)

### Parameter-efficient fine-tuning (PEFT)

#### T

Fine-tuning techniques that modify only a subset of parameters, reducing computational cost. PEFT techniques like LoRA produce smaller artifacts than full fine-tuning, but ownership and portability of these artifacts still requires contractual clarity.

See: [Adapter](#); [Fine-tuning](#); [LoRA \(Low-Rank Adaptation\)](#)

### Patent

#### L

A form of intellectual property that can protect inventions meeting statutory requirements (e.g., novelty, non-obviousness, utility), subject to jurisdiction-specific eligibility rules. In AI contexts, patents may cover model architectures, training techniques, and system implementations.

See: [Intellectual property](#); [Prior art](#); [Trade secret](#)

### PEFT

See: [Parameter-efficient fine-tuning](#)

### Performance

How well a model or system meets task objectives and operational requirements (accuracy, latency, robustness, cost, safety). In contracting, performance is often

expressed as SLAs/SLOs, acceptance criteria, and evaluation benchmarks tied to intended use.

See: [Benchmark](#); [Evaluation \(evals\)](#); [Reliability](#); [SLA/SLO](#)

## Perplexity

T

A metric measuring how well a language model predicts text, with lower values indicating better performance. Perplexity is a technical quality metric primarily useful for model comparison rather than legal analysis.

See: [Benchmark](#); [Evaluation \(evals\)](#)

## Personal data

R

Information relating to an identified or identifiable natural person. AI systems processing personal data trigger privacy obligations under GDPR, state laws, and sector regulations; prompts, outputs, and training data may all contain personal data.

See: [Data subject](#); [Privacy](#)

## Personally Identifiable Information (PII)

Information that can identify an individual, with the precise definition varying by jurisdiction. PII in training data, prompts, or outputs creates compliance obligations; detection and filtering controls are important safeguards.

See: [De-identification](#); [Personal data](#); [Privacy](#)

## PET

See: [Privacy-enhancing technology](#)

## Plugin

A software extension enabling additional functionality in AI systems, such as web browsing, code execution, or database access. Plugins expand AI system capabilities and risks by accessing external resources, and may require separate permissions and security review.

See: [Connector](#); [Tool calling \(function calling\)](#)

## Portability

C

The ability to move an AI workflow, data, or model artifacts from one environment or vendor to another (e.g., switching model providers, migrating vector databases). Portability can turn on data formats, licensing terms, and access to artifacts.

See: [Business continuity](#); [Model artifact](#); [Open weights](#); [Vendor lock-in](#)

## [Post-market monitoring](#)

R

An ongoing obligation under the EU AI Act requiring providers of high-risk AI systems to collect and analyze data on system performance and compliance after deployment. Post-market monitoring is distinct from general operational monitoring; it requires a documented plan proportionate to the system's risks, must feed into the provider's quality management system, and triggers reporting and corrective action obligations when issues are detected.

See: [High-risk AI system](#); [Monitoring](#); [Provider](#); [Quality Management System](#); [Serious incident](#)

## [Post-training](#)

Modifications to a model after pre-training to change behavior or improve usefulness and safety (e.g., instruction tuning, preference optimization, safety tuning, or distillation). Post-training often changes model characteristics and can affect evaluation results, safety properties, and documentation baselines.

See: [Fine-tuning](#); [Pre-training](#)

## [Precision](#)

A metric measuring, of items classified as positive, the fraction that were truly positive. High precision means fewer false positives, which is important for applications where false accusations or unnecessary interventions are costly.

See: [F1 score](#); [False positive](#); [Recall](#)

## [Pre-training](#)

Initial training of a model on large-scale data before any customization or alignment. Pre-training data provenance is the primary driver of copyright and privacy exposure for foundation models.

See: [Foundation model](#); [Post-training](#); [Training](#)

## [Prior art](#)

L

Publicly available information that can be relevant to assessing patentability (or invalidity) of an invention. In AI, prior art may include papers, open source code, model cards, and public model releases.

See: [Open source software](#); [Patent](#); [Publication](#)

## [Privacy](#)

## R

Protection of personal information from unauthorized collection, use, or disclosure. AI creates novel privacy challenges through inference capabilities, behavioral profiling, and potential memorization of training data; multiple legal frameworks apply.

See: [Personal data](#); [Privacy-enhancing technology](#)

### [Privacy by design](#)

## R

Embedding privacy protections into systems from the start rather than as an afterthought. Privacy by design is a GDPR principle requiring consideration of privacy throughout AI development, not just at deployment.

See: [Data minimization](#); [Privacy](#)

### [Privacy-enhancing technology \(PET\)](#)

## R

Technical means of mitigating privacy risks, including encryption, differential privacy, federated learning, and secure enclaves. PETs may support privacy claims but have limitations; understand what protections they actually provide in practice.

See: [Differential privacy](#); [Federated learning](#); [Privacy](#)

### [Probabilistic](#)

A fundamental characteristic of how generative AI models operate: outputs are generated by sampling from learned probability distributions over possible responses rather than by executing logical rules or retrieving stored facts. Even when configured for deterministic operation, a probabilistic model is selecting the statistically most likely output based on training patterns, not computing a provably correct answer. This distinction explains why models can be confidently wrong (hallucination), why explanations of "reasoning" may be post-hoc rationalizations, and why traditional software warranties and performance guarantees require adaptation for AI systems.

See: [Deterministic](#); [Explainability](#); [Hallucination](#); [Neural network](#); [Sampling](#)

### [Prohibited AI practices](#)

## R

AI applications banned under regulatory frameworks such as the EU AI Act. The EU AI Act prohibits certain uses including social scoring, real-time remote biometric identification in public spaces, and emotion recognition in workplaces and schools, with severe penalties for violations.

See: [Biometric data](#); [EU AI Act](#); [High-risk AI system](#)

## Prompt

C T

The input text or instructions provided to an AI system to generate a response. Prompts often contain sensitive business information or personal data; contracts commonly define whether they are logged, retained, or used for improvement.

See: [Input Data](#); [System prompt](#); [User prompt](#)

## Prompt engineering

T

The practice of crafting prompts to achieve desired AI behaviors and outputs. Prompt engineering creates valuable know-how, and prompt libraries may constitute confidential information or trade secrets worth protecting.

See: [Few-shot prompting](#); [Prompting](#); [System prompt](#)

## Prompting

T

The practice of supplying instructions, context, and examples to a model to elicit desired outputs. Prompting includes system/developer prompts and user prompts and is often an alternative to fine-tuning for adapting model behavior.

See: [Few-shot prompting](#); [Prompt engineering](#); [System prompt](#)

## Prompt injection

C R T

An attack in which malicious input causes a model or agent to ignore intended instructions or perform unintended actions (e.g., by overriding system/developer prompts or by exploiting tool integrations). Prompt injection can occur directly via user input or indirectly via retrieved content and is treated as a security risk in many threat models.

See: [Adversarial attack](#); [Jailbreak](#); [Security](#)

## Prompt leakage

T

Unintended disclosure of system prompts or instructions through model outputs. System prompts may contain confidential business logic, competitive information, or security controls; leakage can expose sensitive operational details.

See: [Confidential information](#); [Extraction attack](#); [System prompt](#)

## Prompt library

C T



A collection of tested prompts for specific tasks or use cases. Prompt libraries represent operational controls and potential trade secrets that is often treated as governed assets with change control and access restrictions.

See: [Change control](#); [Confidential information](#); [Prompt engineering](#)

## Provider

R

A party that develops or places an AI system on the market, as defined in the EU AI Act. Provider vs. deployer role allocation determines documentation, conformity assessment, and incident reporting duties under the regulation.

See: [AI system](#); [Deployer](#); [EU AI Act](#)

## Pseudonymization

R

Replacing direct identifiers with tokens while retaining the ability to re-link data under safeguards. Pseudonymized data is generally still personal data under GDPR and similar regimes; do not treat it as equivalent to anonymization.

See: [De-identification](#); [Personal data](#); [Privacy](#)

## Publication

Making information publicly available (e.g., papers, open source repositories, model releases, technical reports). Publications can be relevant to prior art, marketing claims, and transparency commitments.

See: [Model card](#); [Open source software](#); [Prior art](#)

## Purpose limitation

R

A privacy principle requiring data be used only for specified, explicit, and legitimate purposes. Purpose limitation applies when vendors want to reuse prompts, outputs, or logs for improvement; permitted purposes is often defined explicitly.

See: [Data minimization](#); [DPA](#); [Service improvement](#)

# Q

## Quality Assurance (QA)

Processes ensuring AI systems meet quality standards before and during deployment. QA processes support reliability claims and reasonable care arguments; document testing methodologies and acceptance criteria.

See: [Acceptance criteria](#); [Evaluation \(evals\)](#); [Testing](#)

## Quality Management System (QMS)

R

A documented system of policies, procedures, and processes required of providers under the EU AI Act to ensure consistent compliance with regulatory requirements throughout AI system development and operation. QMS obligations for high-risk AI systems include risk management procedures, data governance, technical documentation practices, post-market monitoring, vulnerability identification, incident reporting protocols, and recordkeeping.

See: [AI governance](#); [High-risk AI system](#); [ISO/IEC 42001](#); [Post-market monitoring](#); [Provider](#); [Vulnerability / CVE](#)

## Quantization

C T

Reducing numerical precision of model weights to decrease memory requirements and speed inference. Quantization can change model behavior in subtle ways; treat quantized models as different versions requiring separate validation for regulated deployments.

See: [Edge deployment](#); [Model compression](#); [Performance](#)

# R

## Rate limiting

C

Controls capping API usage to manage load, prevent abuse, and control costs. Rate limits affect service usability, are often disclosed in SLAs, and also serve as security controls against model extraction and denial-of-service attacks.

See: [Abuse](#); [Service Level Agreement / Service Level Objective](#)

## ReAct (Reason \+ Act)

An agent prompting pattern combining reasoning traces with action execution in an interleaved manner. ReAct is a common pattern in agentic AI systems that makes agent decision-making somewhat more transparent.

See: [Agentic AI](#); [Chain of Thought \(CoT\)](#); [Tool calling \(function calling\)](#)

## Reasoning model

C

An AI model specifically trained for complex reasoning, often with extended "thinking" time before responding. Reasoning models like o1 and DeepSeek-R1 represent a capability advancement for analysis tasks but have different cost, latency, and transparency profiles.

See: [Chain of Thought \(CoT\)](#); [Inference-time scaling](#); [Large Reasoning Model](#)

## Recall

A metric measuring, of truly positive items, the fraction that were correctly identified. High recall is critical for safety applications where missing positives is costly, such as fraud detection or medical screening.

See: [F1 score](#); [False negative](#); [Precision](#)

## Records of Processing Activities (ROPA)

### R

Documentation required under GDPR Article 30 cataloging an organization's personal data processing activities, including purposes, data categories, recipients, transfers, retention periods, and security measures. AI systems should be reflected in ROPA entries, with particular attention to training data processing, inference logging, and any cross-border transfers to model providers.

See: [Controller / processor](#); [Cross-border data transfer](#); [Data retention](#); [Personal data](#)

## Records retention

### C

Policies governing how long business records are kept before deletion. AI logs may become business records subject to retention requirements, litigation holds, and deletion rights; align retention periods with legal obligations.

See: [Data retention](#); [E-discovery](#); [Logging](#)

## Recovery Point Objective / Recovery Time Objective (RPO / RTO)

### C

RPO specifies the maximum acceptable data loss measured in time (e.g., "no more than 4 hours of data"); RTO specifies the maximum acceptable downtime before service restoration. These metrics are standard in disaster recovery planning but require special attention for AI systems. Contracts should specify RPO/RTO for each critical component (models, vector stores, knowledge bases, configuration) and clarify whether RTO includes re-indexing and validation time, not just data restoration.

See: [Availability](#); [Business continuity](#); [Service Level Agreement / Service Level Objective](#); [Vector database](#)

## Red teaming

## C R

Adversarial testing by dedicated teams attempting to find system vulnerabilities and failure modes. Red teaming is referenced in the AI Executive Order and EU AI Act; it supports safety claims and regulatory compliance when properly documented.

See: [Adversarial attack](#); [Evaluation \(evals\)](#); [Safety evaluation](#)

## Reinforcement Learning (RL)

A training approach where the model learns by receiving rewards or penalties for its outputs rather than by studying labeled examples. A primary benefit of RL is that in some circumstances it allows synthetic data or self-play to be used in the place of human-labeled data. RL is the foundation for RLHF and is used in game-playing AI and robotics; understanding RL helps explain how models learn to follow instructions.

See: [Reward function](#); [Training](#)

## Reinforcement Learning from AI Feedback (RLAIF)

A variant of RLHF using AI-generated feedback rather than human feedback to reduce cost and scale. RLAIF may have different alignment properties than human feedback and is often combined with Constitutional AI approaches.

See: [Alignment](#); [Constitutional AI](#)

## Reinforcement Learning from Human Feedback (RLHF)

## C T

Training models using human preference data to improve alignment with human values and instructions. RLHF is the dominant technique for aligning LLMs and is key to understanding how models are made to follow instructions and refuse harmful requests.

See: [Alignment](#); [Post-training](#); [Reward model](#)

## Reinforcement Learning with Verifiable Rewards (RLVR)

Training using automatically verifiable outcomes such as correct code execution or valid mathematical proofs. RLVR is used for training reasoning models on tasks with objectively checkable answers.

See: [Reasoning model](#); [Reinforcement Learning](#)

## Reliability

## C R

Consistency and dependability of system behavior over time, including correctness under expected conditions and stability across updates. Reliability is influenced by drift, stochasticity, evaluation coverage, and operational monitoring.

See: [Evaluation \(evals\)](#); [Model drift](#); [Monitoring](#); [Temperature](#)

## [Representation learning](#)

Learning useful internal representations of data that can transfer to multiple downstream tasks. Foundation models learn representations that encode semantic meaning, enabling capabilities like semantic search and few-shot learning.

See: [Embedding](#); [Foundation model](#); [Transfer learning](#)

## [Reproducibility](#)

C

The ability to obtain consistent results from the same inputs across runs. Reproducibility is important for testing, audit, and debugging; non-determinism can be controlled through temperature and seed settings but may affect output quality.

See: [Evaluation \(evals\)](#); [Non-deterministic](#); [Temperature](#)

## [Re-ranking](#)

A retrieval step (often in RAG) that reorders candidate results using a second model (e.g., a cross-encoder or an LLM) to improve relevance. Re-ranking can affect what content is presented to the generation model and therefore affects grounding and auditability.

See: [Answer relevance](#); [Retrieval](#)

## [Resources](#)

In the CORE framework, external assets that an AI system accesses but does not control, such as data, third-party APIs, file systems, knowledge bases, and external services. Resources exist outside the system boundary but are invoked during execution. Resource mapping is essential for data sovereignty compliance, confidentiality protection, and contractual obligation tracking.

See: [Components](#); [Connector](#); [CORE](#); [Data residency](#); [Operations](#); [Tool calling \(function calling\)](#)

## [Responsible AI](#)

R

Principles and practices for developing and deploying AI ethically, safely, and in accordance with human values. Responsible AI frameworks inform governance programs, procurement criteria, and regulatory expectations.

See: [AI governance](#); [AI safety](#); [Ethical AI](#)

## [Retention](#)

Policies and technical practices governing how long data and logs are kept and when they are deleted. In AI systems, retention can apply to prompts, outputs, embeddings, retrieval indexes, and tool call records.

See: [Data governance](#); [Deletion](#); [Litigation hold](#); [Logging](#); [Zero Data Retention](#)

## Retrieval

T

The process of selecting and returning relevant information from a corpus or database (often using keyword or semantic search) to support tasks such as RAG. Retrieval quality affects grounding, hallucination rates, and completeness.

See: [Re-ranking](#); [Semantic search](#); [Vector database](#)

## Retrieval-Augmented Generation (RAG)

C T

An architecture where the system retrieves relevant documents and uses them as context for generation, grounding outputs in specific sources. RAG improves accuracy and reduces hallucination but introduces document access logging, confidentiality, and injection risks that require governance.

See: [Grounding](#); [Knowledge base](#); [Vector database](#)

## Reverse engineering

C

Analyzing a system to understand its design, extract components, or replicate functionality. Reverse engineering restrictions commonly appear in AI terms of service; they may conflict with legitimate security research and interoperability needs.

See: [AUP](#); [Distillation](#); [Model extraction](#); [Trade secret](#)

## Reward function

In reinforcement learning, the function that assigns a numeric score (“reward”) to behaviors, guiding the model toward preferred outcomes. Reward function design influences aligned behavior and can encode tradeoffs.

See: [Alignment](#); [Reinforcement Learning](#)

## Reward model

A model trained to predict human preferences, used to guide RL training by scoring candidate outputs. Reward model quality directly affects alignment effectiveness and the behaviors the final model learns to exhibit.

See: [Alignment](#)

## Right of Publicity

L R

A set of state-law (and sometimes statutory) rights controlling commercial use of a person's name, image, likeness, or voice. In AI contexts, right-of-publicity issues can arise with voice cloning, deepfakes, and digital replicas.

See: [Consent](#); [Deepfake](#); [Digital Replica](#)

## Right to be Forgotten

R

A term commonly referring to deletion rights under certain privacy frameworks (notably in EU law), including circumstances where individuals can seek removal of personal data from search results or other systems. How it applies to AI training data and model artifacts depends on the legal framework and technical implementation.

See: [Deletion](#); [Machine Unlearning](#)

## Right to explanation

R

A data subject's right to understand the logic of automated decisions affecting them. GDPR Article 22 and various U.S. state laws provide explanation rights for significant automated decisions, though the required depth of explanation remains debated.

See: [Automated decision-making](#); [Explainability](#)

## Risk assessment

C R

Systematic evaluation of AI system risks, potential harms, and likelihood of occurrence. Risk assessment is required by the EU AI Act for high-risk systems and expected by governance frameworks; it is often documented and periodically updated.

See: [AI governance](#); [High-risk AI system](#); [NIST AI RMF \(AI Risk Management Framework\)](#)

## Robustness

C

A system's ability to maintain performance under varying conditions, distribution shifts, and adversarial inputs. Robustness claims often specify what conditions and attack types were tested; general robustness guarantees are difficult to provide.

See: [Adversarial attack](#); [Evaluation \(evals\)](#); [Reliability](#)

## S

### SaaS (Software as a Service)

C

Cloud-hosted software accessed over the internet on a subscription basis. Most AI products are delivered as SaaS, with key issues including data handling practices, logging, subprocessor use, and service level commitments. Warning: a SaaS may not be a “product” for purposes of product liability.

See: [Multi-tenancy](#); [Service Level Agreement / Service Level Objective](#)

### Safe harbor

C R

A legal provision shielding parties from liability or enforcement when they meet specified conditions. In the AI context, safe harbors appear in several forms: some state AI laws (such as Colorado's AI Act) provide safe harbors for organizations following recognized frameworks like the NIST AI RMF; Section 230 provides platform immunity that may apply to certain AI-generated content (although the scope is contested); and contractual safe harbors may limit liability when parties follow agreed procedures. Safe harbor protection typically requires documented, good-faith implementation rather than mere assertion of alignment.

See: [AI governance](#); [ISO/IEC 42001](#); [Limitation of liability](#); [NIST AI RMF \(AI Risk Management Framework\)](#)

### Safety evaluation

C R T

Testing specifically focused on identifying unsafe, harmful, or disallowed behaviors. Safety evaluations support governance, regulatory compliance, and reasonable care arguments; they may be required by regulators for certain AI applications.

See: [Evaluation \(evals\)](#); [Guardrails](#); [Red teaming](#)

### Safety policy

C

Rules defining prohibited behaviors for an AI system, implemented through training, prompts, and filtering. Safety policy affects AUP enforcement, duty-of-care arguments, and user expectations; it is often documented and traceable to implemented controls.

See: [AUP](#); [Guardrails](#); [Moderation](#)

### Sampling

Selecting outputs probabilistically from the model's predicted distribution rather than always choosing the highest-probability option. Sampling contributes to output



variability; deterministic modes may be needed for audit and reproducibility requirements.

See: [Non-deterministic](#); [Temperature](#); [Top-p \(nucleus\) sampling](#)

## [Sandboxing](#)

C

Running code or processes in a constrained environment that limits access to sensitive resources. Sandboxing is a key control for agentic systems executing code, supporting security representations and limiting blast radius of failures.

See: [Least privilege](#); [Security](#); [Tool calling \(function calling\)](#)

## [Scalability](#)

C

The ability to maintain performance and availability as usage grows. Scalability affects SLA reliability and business continuity; capacity constraints during high demand can create significant operational and reputational issues.

See: [Performance](#); [Rate limiting](#); [Service Level Agreement / Service Level Objective](#)

## [Scaling law](#)

Empirical relationships showing that model performance improves predictably with increased size, data, and compute. Scaling laws drive investment in larger models and help explain capability improvements, though they don't guarantee specific abilities.

See: [Compute](#); [Foundation model](#); [Parameter](#)

## [Secrets exposure](#)

Unauthorized disclosure of secrets (API keys, credentials, tokens, confidential prompts) through prompts, logs, tool outputs, or model behavior. This is discussed in AI security, incident response, and vendor controls.

See: [Access control](#); [Logging](#); [Prompt injection](#); [Security](#)

## [Secure development lifecycle \(SDLC\)](#)

C

Processes integrating security practices throughout software development. SDLC commitments may appear in security addenda and support reasonable-security arguments in the event of incidents.

See: [Development practices](#); [Security](#)

## [Security](#)

C

Protection of systems and data from unauthorized access, use, disclosure, or destruction. AI systems require security controls appropriate to the sensitivity of data processed and decisions made; common frameworks like SOC 2 provide baseline standards.

See: [Access control](#); [Encryption](#); [SOC 2](#)

## [Security addendum](#)

C

A contract attachment specifying security requirements, controls, and audit rights. AI security addenda often address model-specific concerns including prompt injection defenses, data isolation, logging access, and incident notification.

See: [Contract](#); [DPA](#); [Security](#)

## [Self-attention](#)

T

The attention mechanism applied within a single sequence to capture relationships between different positions. Self-attention is the core mechanism enabling transformers to understand context and relationships in language.

See: [Attention mechanism](#); [Transformer](#)

## [Semantic search](#)

T

Search based on meaning and intent rather than exact keyword matching, typically using embeddings to find semantically similar content. Semantic search powers RAG retrieval; understanding its behavior helps assess retrieval quality and potential failure modes.

See: [Embedding](#); [Vector database](#)

## [Sensitive Personal Information \(SPI\)](#)

R

Categories of personal data treated as especially sensitive under some privacy regimes (e.g., precise geolocation, health data, biometric identifiers, financial data, or data about minors). In AI systems, SPI handling often drives additional controls for collection, processing, retention, access, and disclosures.

See: [Biometric data](#); [Personal data](#); [Privacy](#)

## [Serious incident](#)

R

Under the EU AI Act, an incident or malfunction of a high-risk AI system that directly or indirectly causes death, serious damage to health or property, or serious and irreversible disruption of critical infrastructure. Serious incidents trigger mandatory reporting to market surveillance authorities. Incident classification and reporting procedures should be integrated into the provider's QMS.

See: [EU AI Act](#); [High-risk AI system](#); [Incident response](#); [Post-market monitoring](#); [Provider](#); [Quality Management System](#)

## [Service credits](#)

C

Contractual remedy providing credits against future fees when a vendor fails to meet SLA commitments, typically calculated as a percentage of monthly fees based on the severity and duration of the failure. Service credits are the standard remedy in SaaS agreements.

See: [Availability](#); [Contract](#); [Limitation of liability](#); [Service Level Agreement / Service Level Objective](#)

## [Service improvement](#)

C

Using customer data, including prompts, outputs, and feedback, to improve AI services. Service improvement rights are contentious in negotiations; clearly define whether and how customer content can be used, and provide opt-out mechanisms.

See: [No training on our data](#); [Training](#); [Usage data / telemetry](#)

## [Service Level Agreement / Service Level Objective \(SLA/SLO\)](#)

C

Contractual commitments on service performance metrics such as availability, latency, and throughput. AI SLAs often address model-specific issues including response time variability, rate limits, and model deprecation notice.

See: [Availability](#); [Latency](#); [Uptime](#)

## [Shadow AI](#)

Unauthorized use of AI tools by employees outside official channels and governance processes. Shadow AI increases confidentiality breach risk, privilege waiver concerns, and compliance exposure; it requires both policy controls and technical measures to address.

See: [Acceptable use](#); [AI governance](#)

## [Similarity search](#)

T

Finding items semantically close to a query by comparing embeddings in vector space. Similarity thresholds affect retrieval quality and is often documented for regulated deployments where retrieval completeness matters.

See: [Cosine similarity](#); [Vector database](#)

## [SLA/SLO](#)

See: [Service Level Agreement / Service Level Objective](#)

## [Slop](#)

Pejorative term for low-quality, mass-produced AI-generated content, typically created with minimal human oversight or input.

See: [AUP](#); [Content provenance](#); [Generated content](#); [Generative AI \(GenAI\)](#); [Synthetic media](#); [Watermarking](#)

## [Small Language Model \(SLM\)](#)

C T

A language model with fewer parameters (typically  $<10B$ ) designed for efficiency, lower latency, and often local/edge deployment. Unlike larger models, SLMs can often run on consumer hardware (laptops, phones) without sending data to a cloud provider, changing the privacy and security risk profile.

See: [Distillation](#); [Edge deployment](#); [Quantization](#)

## [SOC 2](#)

C

An audit framework and report on controls for security, availability, processing integrity, confidentiality, and privacy. SOC 2 reports are standard due diligence for AI vendors; confirm the report scope includes AI-relevant systems and controls.

See: [Audit](#); [Security](#)

## [Software Bill of Materials \(SBOM\)](#)

C

An inventory of software components, dependencies, and their versions used in a system. SBOMs support open source license compliance, vulnerability management, and supply chain security; they are increasingly requested in enterprise procurement.

See: [License compliance](#); [Open source](#); [Supply chain security](#)

## [Source attribution](#)

Linking generated outputs to the documents or data used to produce them. Attribution supports defensibility and user trust but can be incorrect or fabricated; validate citation mechanisms through testing.

See: [Citation](#); [Grounding](#)

## [SPI](#)

See: [Sensitive Personal Information](#)

## [Statement of Work \(SOW\)](#)

### C

A contract section or attachment describing specific services, deliverables, and acceptance criteria. AI SOWs often specify model boundaries, evaluation metrics, change control procedures, and artifact ownership.

See: [Acceptance criteria](#); [Contract](#); [Deliverables](#)

## [State Space Model \(SSM\)](#)

### T

A neural architecture modeling sequences using continuous state representations, offering efficient processing of long sequences. SSMs like Mamba are alternatives to transformers for applications requiring very long context windows or efficient inference.

See: [Context window](#); [Mamba](#); [Transformer](#)

## [Structured output](#)

### C

Constraining AI outputs to a defined schema such as JSON, XML, or specific formats for reliable parsing. Structured output improves auditability, reduces parsing errors, and enables automated processing in enterprise workflows.

See: [Function calling](#); [JSON mode](#)

## [Subprocessor](#)

### C

A third party engaged by a vendor to process data on its behalf. Subprocessor lists, notification of changes, and flow-down of data protection obligations are central provisions in DPAs governing AI services.

See: [Cloud provider](#); [Controller / processor](#); [DPA](#)

## [Substantial modification](#)

### R

Under the EU AI Act, a change to a high-risk AI system after initial placing on the market that affects compliance with regulatory requirements or alters the system's intended purpose. Changes to training data, model architecture, or safety controls may constitute substantial modifications even if commercial framing remains unchanged.

See: [Change control](#); [Deployer](#); [EU AI Act](#); [High-risk AI system](#); [Model update](#); [Provider](#)

### [Substantial similarity](#)

L

The legal standard for copyright infringement based on whether works are sufficiently similar in protected expression. Substantial similarity analysis is central to AI copyright litigation.

See: [Copyright](#); [Fair use](#)

### [Summarization](#)

Condensing longer content into shorter form while preserving key information. AI summarization may miss important details, introduce errors, or change emphasis; users often verify accuracy for legal and business-critical matters.

See: [Hallucination](#); [Truncation](#)

### [Supervised learning](#)

Machine learning from labeled examples that map inputs to correct outputs. Supervised learning requires labeled training data, raising data rights, annotation labor, and quality control considerations.

See: [Labeling](#); [Machine learning](#); [Training](#)

### [Supply chain security](#)

C

Protecting against risks from third-party components, data sources, and service providers. AI supply chains include foundation models, training datasets, open source libraries and other components, as well as third-party cloud infrastructure; each introduces potential vulnerabilities.

See: [Data poisoning](#); [Open source](#)

### [Sustainability](#)

Environmental and resource considerations of AI development and deployment (energy use, water use, hardware lifecycle). Sustainability may be discussed in procurement, ESG reporting, and policy debates about compute-intensive systems.

See: [Compute](#); [Datacenter](#)

## Synthetic data

Data generated by a computer process (possibly an AI system) rather than collected from real-world sources. Synthetic data can reduce privacy concerns and augment training sets but may not accurately represent real-world distributions or edge cases.

See: [Data augmentation](#); [Model collapse](#); [Privacy](#); [Training data](#)

## Synthetic media

AI-generated content including images, audio, video, and text created to resemble authentic content. Synthetic media raises authenticity, deepfake, evidence authentication, and misinformation concerns; provenance and detection tools are evolving. A number of laws require the disclosure of synthetic media.

See: [Content provenance](#); [Deepfake](#); [Generative AI \(GenAI\)](#)

## System card

Documentation describing a complete AI system's design, intended use, capabilities, limitations, and safety measures. System cards provide more complete context than model cards by covering the full deployed system, not just the underlying model.

See: [Architecture](#); [CORE](#); [Documentation](#); [Model card](#); [Technical documentation](#)

## Systemic risk

R

Risk that AI system failures or misuse could have wide-ranging negative impacts on society, the economy, or critical infrastructure. The EU AI Act imposes additional transparency and evaluation obligations on GPAI models posing systemic risk.

See: [EU AI Act](#); [GPAI](#); [High-risk AI system](#)

## System prompt

C T

Instructions provided to the model that frame its role, capabilities, and constraints, typically hidden from end users. System prompts contain business logic and operational controls that may be confidential; prompt leakage is a security concern.

See: [Confidential information](#); [Prompt](#); [Prompt leakage](#)

## T

## Technical documentation

R

Formal documentation of AI system design, development, testing, and deployment required by regulatory frameworks. The EU AI Act requires comprehensive technical documentation for high-risk systems; retain documentation for compliance and liability defense.

See: [Conformity assessment](#); [Model card](#); [System card](#)

## Temperature

C T

A parameter controlling output randomness, where lower values produce more deterministic and focused outputs. Temperature settings affect reproducibility, creativity, and consistency; document settings for regulated or audited workflows.

See: [Hyperparameter](#); [Non-deterministic](#); [Sampling](#)

## Tensor Processing Unit (TPU)

T

Google's custom-designed AI accelerator hardware optimized for neural network operations. TPUs are alternatives to GPUs with different availability, pricing, and vendor dependencies.

See: [Compute](#)

## Terms of Service

C

Contract terms governing use of a product or service (often online), including license grants, restrictions, disclaimers, limitation of liability, and incorporated policies such as an AUP. Terms of Service are commonly relevant to authorization, enforcement, and dispute resolution.

See: [Acceptable Use Policy](#); [Contract](#); [Limitation of liability](#); [Warranty](#)

## Testing

Methods for assessing whether a model or system meets requirements and behaves safely and reliably (e.g., unit tests, red teaming, evals, regression tests). Testing results are often used in acceptance, governance, and incident response.

See: [Acceptance criteria](#); [Change control](#); [Evaluation \(evals\)](#); [Red teaming](#)

## Text classification

A common inference task where a model assigns one or more labels to text (e.g., spam/not spam, topic tagging, sentiment). Classification models may be built with traditional ML, deep learning, or LLM prompting.

See: [Classification](#); [Inference](#); [NLP](#); [Supervised learning](#)



## Text-to-image

AI generation of images from text descriptions or prompts. Text-to-image systems raise copyright questions about training data and outputs, trademark concerns, and deepfake risks.

See: [Diffusion model](#); [Generative AI \(GenAI\)](#); [Image generation](#)

## Threat model

An analysis identifying potential attackers, attack vectors, assets at risk, and security boundaries for a system. Threat models inform security requirements, testing priorities, and control design; they should be documented and updated as systems evolve.

See: [Adversarial attack](#); [OWASP Top 10 for LLMs](#); [Prompt injection](#); [Risk assessment](#); [Security](#)

## Throughput

C T

The rate of processing, typically measured in requests per second or tokens per second. Throughput affects system capacity and cost; it may be specified in SLAs and is often tested under realistic load conditions.

See: [Latency](#); [Rate limiting](#); [Service Level Agreement / Service Level Objective](#)

## Token

C T

The unit of text that models actually process. Tokens are not words: common words may be single tokens, but less common words are split into pieces ("contract" might be one token; "indemnification" might be three). A rough approximation is 0.75 words per token for English. Token counts determine API costs, which are typically priced per token, and they define context window limits. Understanding tokenization helps explain why non-English text and technical terminology often perform worse—they require more tokens to represent the same meaning.

See: [Context window](#); [Rate limiting](#); [Tokenization](#)

## Tokenization

T

The process of converting text into tokens that the model can process. Tokenization affects how different languages, technical content, and special characters are handled, potentially causing issues with non-English text or domain-specific terminology. In contrast to vectorization, tokenization is a reversible, direct translation of the input.

See: [Context window](#); [Token](#); [Vectorization](#)

## Tool calling (function calling)

C T

A mechanism where the model outputs structured instructions to invoke external functions, APIs, or services. Tool calling is a key inflection point for risk because it enables AI systems to take real-world actions; permissions, logging, and approval workflows are essential controls.

See: [Agentic AI](#); [Least privilege](#)

## Tool output injection

A risk where untrusted tool outputs (e.g., web content, search results, emails) introduce malicious instructions or data that affect subsequent model behavior. This is closely related to indirect prompt injection in tool-enabled systems.

See: [Indirect prompt injection](#); [Prompt injection](#); [Tool calling \(function calling\)](#)

## Tool permissions

C

Controls specifying what actions and resources AI agents can access. Permissions are primary controls for limiting potential harm from agentic systems; implement least privilege principles and require explicit authorization for sensitive operations.

See: [Agentic AI](#); [Excessive agency](#); [Least privilege](#)

## Top-k sampling

T

A sampling method limiting next-token selection to the k highest-probability options. Top-k is one of several sampling parameters affecting output determinism; document settings for reproducibility requirements.

See: [Sampling](#); [Temperature](#); [Top-p \(nucleus\) sampling](#)

## Top-p (nucleus) sampling

T

A sampling method limiting selection to tokens within a cumulative probability threshold. Like temperature, top-p affects output variability and is often documented for audit and reproducibility purposes.

See: [Sampling](#); [Temperature](#); [Top-k sampling](#)

## Trade secret

C

Information deriving economic value from not being generally known and subject to reasonable secrecy efforts. Model weights, training data compositions, system prompts,

and prompt libraries may qualify as trade secrets; logging and vendor access can undermine secrecy claims.

See: [Confidential information](#); [IP](#); [Weights](#)

## Training

C

The process of adjusting model weights using data to improve performance on target tasks. Training definitions are central to negotiations about customer data use, distinguishing training from inference, evaluation, and service improvement.

See: [Fine-tuning](#); [Post-training](#); [Pre-training](#)

## Training data

C L

Data used to train or fine-tune a model, including text, images, code, and other content. Training data provenance drives copyright and privacy exposure for AI systems; it is central to indemnity negotiations and ongoing litigation.

See: [Copyright](#); [Dataset documentation](#); [IP indemnity](#)

## Transfer learning

Adapting a pre-trained model to new tasks rather than training from scratch, leveraging learned representations. Transfer learning underlies most commercial AI applications and drives questions about base model rights versus adaptation rights.

See: [Fine-tuning](#); [Foundation model](#); [Pre-training](#)

## Transformative use

L

In copyright fair use analysis, whether a use adds new meaning, message, or purpose to the original work. Transformative use is a key factor in AI training data litigation; courts are actively deciding how the doctrine applies to machine learning.

See: [Copyright](#); [Fair use](#); [Training data](#)

## Transformer

C T

The neural network architecture underlying modern LLMs and most other frontier AI systems. Transformers process input by converting it to tokens, then using attention mechanisms to determine which parts of the input are relevant to each other. This architecture enables models to handle long-range dependencies in text, such as understanding that a pronoun in one sentence refers to a noun several paragraphs

earlier. "Transformer-based" in vendor materials signals a model with LLM-like capabilities.

See: [Architecture](#); [Attention mechanism](#); [LLM](#)

## Transparency

R L

The degree to which information about an AI system is disclosed and understandable (e.g., intended use, data sources at a high level, limitations, evaluation results, safety controls). Transparency is commonly discussed in policy, procurement, and consumer protection contexts.

See: [Disclosure](#); [Documentation](#); [Model card](#); [System card](#)

## Tree-of-thoughts prompting

T

A prompting framework that explores multiple reasoning paths in a tree structure rather than a single chain. Tree-of-thoughts is an advanced technique for complex reasoning tasks requiring evaluation of alternatives.

See: [Chain of Thought \(CoT\)](#); [Prompting](#); [Reasoning model](#)

## Truncation

Dropping content when inputs exceed context window limits or other constraints, often without explicit notification. Truncation can undermine reliability because users may not know the system ignored portions of their input; it creates risk for legal document review.

See: [Context window](#); [Summarization](#); [Token](#)

## Trusted execution environment

A hardware- or enclave-based environment designed to protect code and data during execution (e.g., isolating workloads from a host OS). TEEs are discussed in confidential computing and can be used to reduce exposure when processing sensitive data.

See: [Confidential computing](#); [Encryption](#); [Security](#)

## Trustworthy AI

R

AI systems exhibiting characteristics including validity, reliability, safety, security, transparency, fairness, and privacy protection. The NIST AI RMF defines trustworthiness characteristics that inform governance programs and regulatory expectations.

See: [AI governance](#); [NIST AI RMF \(AI Risk Management Framework\)](#); [Responsible AI](#)

## U

### Unbounded consumption

#### R

A vulnerability where AI systems consume excessive computational or financial resources without limits. OWASP Top 10 for LLMs identifies this as a key risk; rate limiting, budget controls, and resource monitoring are mitigations.

See: [OWASP Top 10 for LLMs](#); [Rate limiting](#); [Security](#)

### Unsupervised learning

Learning patterns from unlabeled data without explicit correct answers. Unsupervised techniques like clustering can create sensitive inferences about individuals or groups even without labeled protected attributes.

See: [Clustering](#); [Embedding](#); [Machine learning](#)

### Uptime

#### C

The percentage of time a service is operational and accessible over a measurement period. Uptime commitments in AI SLAs often address model-specific failure modes; pair with incident response and communication terms.

See: [Availability](#); [Service Level Agreement / Service Level Objective](#)

### Usage data / telemetry

#### C

Operational data about service use, which may include prompts, outputs, timestamps, and performance metrics. Telemetry is often where service improvement occurs; contracts commonly define permitted collection, retention periods, and use restrictions.

See: [Data retention](#); [Logging](#); [Service improvement](#)

### User prompt

#### T

The end-user's input to an AI system, as distinguished from system prompts set by developers. User prompts frequently contain personal or confidential information; define logging practices and provide appropriate privacy notices. In many cases, the user prompt is the proximate cause of the AI system output.

See: [Input Data](#); [Logging](#); [Privacy](#)

## V

### Validation

Testing that a system meets its specified requirements and performs acceptably for intended uses. Validation is key to acceptance criteria and regulatory compliance; document validation methodology and results.

See: [Acceptance criteria](#); [Evaluation \(evals\)](#); [Testing](#)

### Value alignment

Ensuring AI systems pursue goals and exhibit behaviors consistent with human values and intentions. Value alignment is a central AI safety concept; misalignment between system objectives and human values can cause harmful behavior even without adversarial attack.

See: [AI safety](#); [Alignment](#)

### Variational Autoencoder (VAE)

T

A generative model architecture that encodes data as probability distributions in latent space rather than fixed points. VAEs are used in some image generation systems and for learning compressed representations.

See: [Autoencoder](#); [Generative AI \(GenAI\)](#); [Latent space](#)

### Vector

C T

A list of numbers representing an item in a mathematical space. Embeddings are not a direct translation of the item content; instead, vectors are a one-way transformation that encodes semantic meaning. Nevertheless, vectors derived from sensitive content may themselves be sensitive; assess whether embeddings constitute personal data or confidential information.

See: [Embedding](#); [Latent space](#)

### Vector database

C T

A database optimized for storing embeddings and performing fast similarity search at scale. Vector databases often store representations of sensitive enterprise content; access control, encryption, and retention policies are critical security considerations.

See: [Embedding](#); [Semantic search](#)

### Vectorization

## T

Converting inputs into numeric vector representations for model processing by representing their statistical values in a virtual high-dimensional space. Vectorization raises questions about whether derived representations retain the legal significance of source content, including personal data characteristics.

See: [Embedding](#); [Tokenization](#); [Vector](#)

## Vendor lock-in

### C

Dependence on a specific vendor that makes switching costly or difficult. AI lock-in can arise from proprietary formats, fine-tuned models, prompt libraries, and integrated workflows; consider portability of all assets when selecting vendors.

See: [Business continuity](#); [Portability](#); [Version pinning](#)

## Verification

Confirming that outputs, claims, or system behaviors are accurate and meet requirements. Human verification is a key control for managing hallucination risk in high-stakes applications; define verification requirements and responsibilities.

See: [Hallucination](#); [Human-in-the-loop](#)

## Version pinning

### C

Locking to a specific model version to prevent unplanned behavior changes from updates. Pinning is important for validated and regulated workflows; contracts commonly address version availability, deprecation notice, and migration support.

See: [Change control](#); [Model drift](#); [SLA/SLO](#)

## Vision-language model

### T

A multimodal model that processes both visual inputs (images/video) and text, enabling tasks such as captioning, visual question answering, and document understanding.

See: [Computer vision](#); [Multimodal model](#)

## Vulnerability / CVE

### C

A weakness in software, systems, or models that could be exploited to cause harm; CVE (Common Vulnerabilities and Exposures) is the standardized system for identifying and cataloging such weaknesses. Traditional software vulnerabilities in AI systems or components follow established CVE disclosure and patching processes—and these

processes may be required for some systems due to regulations like the EU AI Act and Cyber Resilience Act (CRA). AI-specific vulnerabilities like prompt injection and jailbreaks are less standardized; MITRE ATLAS and OWASP Top 10 for LLMs provide emerging taxonomies but lack CVE-style universal identifiers. Contracts should address vulnerability disclosure timelines, patching obligations, and notification requirements for both software and model-level vulnerabilities. When evaluating vendor security posture, assess both traditional vulnerability management (CVE monitoring, patch cadence) and AI-specific security practices.

See: [Cyber Resilience Act](#); [Incident response](#); [OWASP Top 10 for LLMs](#); [Prompt injection](#); [Security](#)

## W

### Warranty

#### C

A contractual assurance about a product or service (e.g., performance, conformity to documentation, security controls, non-infringement). In AI agreements, warranties are frequently scoped by intended use, evaluation criteria, and exclusions for customer inputs, misuse, or model updates.

See: [Indemnity](#); [Limitation of liability](#); [Service Level Agreement / Service Level Objective](#)

### Watermarking

#### R

Embedding detectable signals in AI-generated content to indicate its origin or enable provenance tracking. Watermarks can support authenticity verification but may be fragile or removable; do not overstate their reliability as detection mechanisms.

See: [Content provenance](#); [Metadata](#)

### Weights

#### C

The learned numerical parameters of a neural network that determine its behavior; the core model artifact. Weights are frequently treated as valuable trade secrets and key IP; access restrictions, licensing terms, and security controls are important diligence topics.

See: [Model](#); [Open weights](#); [Parameter](#)

## X

### XAI



See: [Explainable AI](#)

## Z

### [Zero Data Retention \(ZDR\)](#)

C

A vendor commitment not to retain customer prompts, outputs, or associated data beyond the duration necessary to process the request and return a response. ZDR is often offered as an API option or enterprise tier feature to address confidentiality, privacy, and "no training on our data" concerns. However, ZDR policies vary significantly in scope: clarify whether ZDR covers abuse monitoring logs, trust and safety reviews, debugging data, error logs, metadata, and cached embeddings. Some vendors retain data briefly (e.g., 30 days) for abuse detection even under "zero retention" labels; others exclude certain content flagged for safety review. ZDR does not address what happened to data *before* the policy was enabled, nor does it prevent data exposure during transmission or processing.

See: [Confidential information](#); [Data retention](#); [Logging](#); [No training on our data](#); [Service improvement](#); [Usage data / telemetry](#)

### [Zero-shot prompting](#)

T

Prompting a model to perform a task without providing examples, relying entirely on the model's pre-existing capabilities. Zero-shot performance varies significantly by task; testing often verifies the model can reliably perform intended uses without examples.

See: [Few-shot prompting](#); [In-context learning](#); [Prompting](#)

[Top](#) [A](#) [B](#) [C](#) [D](#) [E](#) [F](#) [G](#) [H](#) [I](#) [J](#) [K](#) [L](#) [M](#) [N](#) [O](#) [P](#) [Q](#) [R](#) [S](#) [T](#) [U](#) [V](#) [W](#) [X](#) [Z](#)